

Model

Eduardo Yuki Yada

Imports

```
library(tidyverse)
library(yaml)
library(tidymodels)
library(usemodels)
library(vip)
```

Loading data

```
load('../dataset/processed_data.RData')
load('../dataset/processed_dictionary.RData')

columns_list <- yaml.load_file("../auxiliar/columns_list.yaml")

outcome_column <- params$outcome_column
features_list <- params$features_list
```

Filtering eligible patients

```
df = df %>%
  filter(disch_outcomes_t0 == 0)

df %>% dim
```

```
## [1] 15766 239
```

Eligible features

```
eligible_columns = df_names %>%
  filter(momento.aquisicao == 'Admissão t0') %>%
  .$variable.name

exception_columns = c('death_intraop', 'death_intraop_1')

correlated_columns = c('year_procedure_1', # com year_adm_t0
  'age_surgery_1', # com age
  'admission_pre_t0_count', # com admission_t0
  'atb', # com meds_antimicrobianos
  'classe_meds_cardio_qtde', # com classe_meds_qtde
  'suporte_hemod' # com proced_invasivos_qtde
)

eligible_features = eligible_columns %>%
  base::intersect(c(columns_list$categorical_columns, columns_list$numerical_columns)) %>%
  setdiff(c(exception_columns, correlated_columns))
```

```

if (is.null(features_list)) {
  features = eligible_features
} else {
  features = base::intersect(eligible_features, features_list)
}

length(features)

## [1] 78
print(features)

## [1] "education_level"      "underlying_heart_disease" "heart_disease"
## [4] "nyha_basal"           "prior_mi"                 "heart_failure"
## [7] "transplant"           "endocardites"             "hemodialysis"
## [10] "comorbidities_count"  "procedure_type_1"         "reop_type_1"
## [13] "procedure_type_new"   "cied_final_1"             "cied_final_group_1"
## [16] "admission_t0"         "admission_pre_t0_180d"    "icu_t0"
## [19] "dialysis_t0"          "disch_outcomes_t0"        "admission_t0_emergency"
## [22] "aco"                  "antiarritmico"            "betabloqueador"
## [25] "ieca_bra"             "dva"                      "digoxina"
## [28] "estatina"             "diuretico"                "vasodilatador"
## [31] "insuf_cardiaca"       "espironolactona"          "bloq_calcio"
## [34] "antiplaquetario_ev"   "insulina"                 "anticonvulsivante"
## [37] "psicofarmacos"        "antifungico"              "antiviral"
## [40] "classe_meds_qtde"     "meds_cardiovasc_qtde"     "meds_antimicrobianos"
## [43] "cec"                  "transplante_cardiaco"     "outros_proced_cirurgicos"
## [46] "icp"                  "intervencao_cv"           "cateterismo"
## [49] "eletrofisiologia"     "cateter_venoso_central"    "proced_invasivos_qtde"
## [52] "cve_desf"             "transfusao"               "equipe_multiprof"
## [55] "ecg"                  "holter"                   "tilt_teste"
## [58] "metodos_graficos_qtde" "laboratorio"              "cultura"
## [61] "analises_clinicas_qtde" "citologia"                "biopsia"
## [64] "histopatologia_qtde"   "angio_rm"                 "angio_tc"
## [67] "cintilografia"         "ecocardiograma"           "endoscopia"
## [70] "flebografia"          "pet_ct"                   "ultrassom"
## [73] "tomografia"           "radiografia"              "ressonancia"
## [76] "exames_imagem_qtde"    "bic"                      "mpp"

```

Train test split (70%/30%)

```

set.seed(42)

df[columns_list$outcome_columns] <- lapply(df[columns_list$outcome_columns], factor)

df_split <- initial_split(df %>% dplyr::select(all_of(c(features, outcome_column))),
                          prop = .7, strata = all_of(outcome_column))
df_train <- training(df_split)
df_test <- testing(df_split)

dim(df_train)[1] / dim(df)[1]

## [1] 0.6999873
dim(df_test)[1] / dim(df)[1]

## [1] 0.3000127

```

Global parameters

```
k = 4 # Number of folds for cross validation

set.seed(234)
df_folds <- vfold_cv(df_train, v = k,
                     strata = all_of(outcome_column))
```

Functions

```
validation = function(model_fit, new_data) {
  library(pROC)
  library(caret)

  test_predictions_prob <-
    predict(model_fit, new_data = new_data, type = "prob") %>%
    rename_at(vars(starts_with(".pred_")), ~ str_remove(., ".pred_")) %>%
    .$`1`

  pROC_obj <- roc(
    new_data[[outcome_column]],
    test_predictions_prob,
    smoothed = TRUE,
    # arguments for ci
    ci = TRUE,
    ci.alpha = 0.9,
    stratified = FALSE,
    # arguments for plot
    plot = TRUE,
    auc.polygon = TRUE,
    max.auc.polygon = TRUE,
    grid = TRUE,
    print.auc = TRUE,
    show.thres = TRUE
  )

  sens.ci <- ci.se(pROC_obj)
  plot(sens.ci, type = "shape", col = "lightblue")
  plot(sens.ci, type = "bars")

  test_predictions_class <-
    predict(model_fit, new_data = new_data, type = "class") %>%
    rename_at(vars(starts_with(".pred_")), ~ str_remove(., ".pred_")) %>%
    .$class

  conf_matrix = table(test_predictions_class, new_data[[outcome_column]])

  confusionMatrix(conf_matrix) %>% print

  return(pROC_obj)
}
```

Boosted Tree (XGBoost)

```
xgboost_recipe <-
  recipe(formula = sprintf("%s ~ .", outcome_column) %>% as.formula, data = df_train) %>%
  step_string2factor(all_nominal_predictors()) %>%
  step_novel(all_nominal_predictors()) %>%
  step_dummy(all_nominal_predictors(), one_hot = TRUE) %>%
```

```

step_zv(all_predictors())

xgboost_spec <- boost_tree(
  trees = 500,
  tree_depth = tune(),
  min_n = tune(),
  loss_reduction = tune(),
  sample_size = tune(),
  mtry = tune(),
  learn_rate = tune()
) %>%
  set_engine("xgboost") %>%
  set_mode("classification")

xgboost_grid <- grid_latin_hypercube(
  tree_depth(),
  min_n(),
  loss_reduction(),
  sample_size = sample_prop(),
  finalize(mtry(), df_train),
  learn_rate(),
  size = 10
)

xgboost_workflow <-
  workflow() %>%
  add_recipe(xgboost_recipe) %>%
  add_model(xgboost_spec)

xgboost_tune <-
  xgboost_workflow %>%
  tune_grid(resamples = df_folds,
            grid = xgboost_grid)

xgboost_tune %>%
  show_best("roc_auc")


## # A tibble: 5 x 12
##   mtry min_n tree_depth learn_rate loss_reduction sample_size .metric .estimator mean
##   <int> <int>    <int>      <dbl>         <dbl>         <dbl> <chr>   <chr>    <dbl>
## 1    19    18        10  0.00197         1.45e- 1         0.748 roc_auc binary  0.683
## 2    79    24         9  0.00120         1.38e- 7         0.942 roc_auc binary  0.677
## 3    40    14         7  0.0000268       1.69e- 9         0.906 roc_auc binary  0.671
## 4    65    26         3  0.000000996     4.99e- 1         0.247 roc_auc binary  0.665
## 5    36    34        15  0.00000849     1.43e-10         0.551 roc_auc binary  0.664
## # ... with 3 more variables: n <int>, std_err <dbl>, .config <chr>

best_xgboost <- xgboost_tune %>%
  select_best("roc_auc")


xgboost_tune %>%
  collect_metrics() %>%
  filter(.metric == "roc_auc") %>%
  select(mean, mtry:sample_size) %>%
  pivot_longer(mtry:sample_size,
               values_to = "value",
               names_to = "parameter"
  ) %>%
  ggplot(aes(value, mean, color = parameter)) +
  geom_point(alpha = 0.8, show.legend = FALSE) +
  facet_wrap(~parameter, scales = "free_x") +

```

```
labs(x = NULL, y = "AUC")
```




```
final_xgboost_workflow <-  
  xgboost_workflow %>%  
  finalize_workflow(best_xgboost)  
  
last_xgboost_fit <-  
  final_xgboost_workflow %>%  
  last_fit(df_split)  
  
final_xgboost_fit <- extract_workflow(last_xgboost_fit)  
  
xgboost_auc = validation(final_xgboost_fit, df_test)
```



```
## Confusion Matrix and Statistics
##
##
## test_predictions_class    0    1
##                0 4540  190
##                1    0    0
##
##              Accuracy : 0.9598
##              95% CI   : (0.9538, 0.9652)
##    No Information Rate : 0.9598
##    P-Value [Acc > NIR] : 0.5193
##
##              Kappa : 0
##
##  McNemar's Test P-Value : <2e-16
##
##              Sensitivity : 1.0000
##              Specificity : 0.0000
##    Pos Pred Value : 0.9598
##    Neg Pred Value :    NaN
##    Prevalence : 0.9598
##    Detection Rate : 0.9598
##    Detection Prevalence : 1.0000
##    Balanced Accuracy : 0.5000
##
##    'Positive' Class : 0
##
```

```
final_xgboost_fit %>%
  fit(data = df_train) %>%
  extract_fit_parsnip() %>%
  vip(geom = "point")
```



GLM

```
glmnet_recipe <-
  recipe(formula = sprintf("%s ~ .", outcome_column) %>% as.formula, data = df_train) %>%
  step_string2factor(all_nominal_predictors()) %>%
  step_novel(all_nominal_predictors()) %>%
  step_dummy(all_nominal_predictors()) %>%
  step_zv(all_predictors()) %>%
  step_normalize(all_numeric_predictors())


glmnet_spec <-
  logistic_reg(penalty = 0) %>%
  set_mode("classification") %>%
  set_engine("glmnet")

glmnet_workflow <-
  workflow() %>%
  add_recipe(glmnet_recipe) %>%
  add_model(glmnet_spec)

glm_fit <- glmnet_workflow %>%
  fit(df_train)

# glm_fit %>%
#   pull_workflow_fit() %>%
#   tidy()

glm_auc = validation(glm_fit, df_test)
```



```
## Confusion Matrix and Statistics
##
##
## test_predictions_class    0    1
##                0 4540  189
##                1    0    1
##
##              Accuracy : 0.96
##              95% CI : (0.9541, 0.9654)
##    No Information Rate : 0.9598
##    P-Value [Acc > NIR] : 0.4898
##
##              Kappa : 0.0101
##
## McNemar's Test P-Value : <2e-16
##
##              Sensitivity : 1.000000
##              Specificity : 0.005263
##    Pos Pred Value : 0.960034
##    Neg Pred Value : 1.000000
##    Prevalence : 0.959831
##    Detection Rate : 0.959831
##    Detection Prevalence : 0.999789
##    Balanced Accuracy : 0.502632
##
##    'Positive' Class : 0
##
```

Decision Tree

```
tree_recipe <-
  recipe(formula = sprintf("%s ~ .", outcome_column) %>% as.formula, data = df_train) %>%
```



```

step_string2factor(all_nominal_predictors()) %>%
step_novel(all_nominal_predictors()) %>%
step_zv(all_predictors())

tree_spec <-
  decision_tree(cost_complexity = tune(),
                tree_depth = tune()) %>%
  set_mode("classification") %>%
  set_engine("rpart")

tree_grid <- grid_latin_hypercube(cost_complexity(),
                                  tree_depth(),
                                  size = 20)


tree_workflow <-
  workflow() %>%
  add_recipe(tree_recipe) %>%
  add_model(tree_spec)

tree_tune <-
  tree_workflow %>%
  tune_grid(resamples = df_folds,
            grid = tree_grid)

tree_tune %>%
  collect_metrics()

autoplot(tree_tune, metric = "roc_auc")

```



```

tree_tune %>%
  show_best("roc_auc")

best_tree <- tree_tune %>%

```

```

select_best("roc_auc")


final_tree_workflow <-
  tree_workflow %>%
  finalize_workflow(best_tree)

last_tree_fit <-
  final_tree_workflow %>%
  last_fit(df_split)

final_tree_fit <- extract_workflow(last_tree_fit)

tree_auc = validation(final_tree_fit, df_test)


```



```

if (tree_auc$auc > 0.55){
  final_tree_fit %>%
    extract_fit_parsnip() %>%
    vip()
}

```



Random Forest

```
rf_recipe <-
  recipe(formula = sprintf("%s ~ .", outcome_column) %>% as.formula, data = df_train) %>%
  step_string2factor(all_nominal_predictors()) %>%
  step_novel(all_nominal_predictors()) %>%
  step_zv(all_predictors()) %>%
  step_dummy(all_nominal_predictors()) %>%
  step_impute_mode(all_nominal_predictors()) %>%
  step_impute_mean(all_numeric_predictors())

rf_spec <-
  rand_forest(mtry = tune(),
              trees = 1000,
              min_n = tune()) %>%
  set_mode("classification") %>%
  set_engine("ranger")


rf_grid <- grid_latin_hypercube(mtry(range = c(1, 10)),
                               min_n(),
                               size = 10)

rf_workflow <-
  workflow() %>%
  add_recipe(rf_recipe) %>%
  add_model(rf_spec)

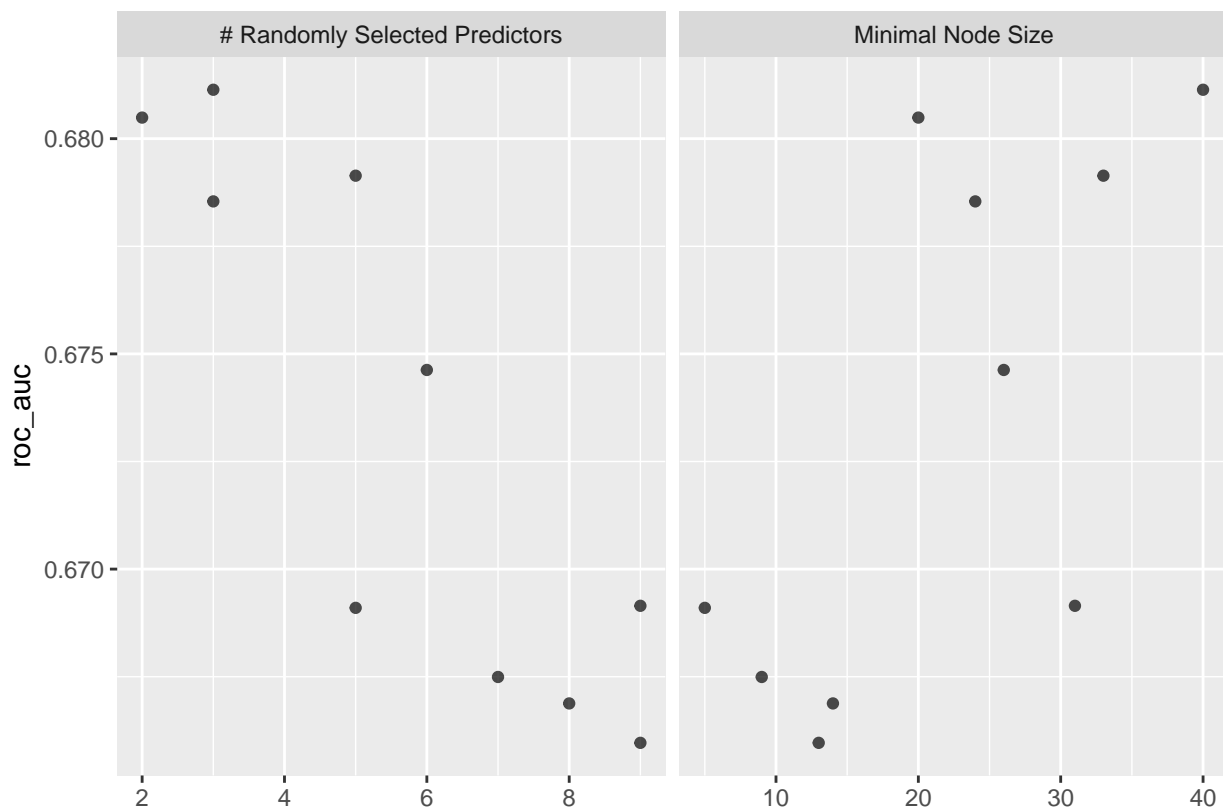
rf_tune <-
  rf_workflow %>%
  tune_grid(resamples = df_folds,
            grid = rf_grid)
```

```
rf_tune %>%
  collect_metrics()

rf_tune %>%
  collect_metrics() %>%
  filter(.metric == "roc_auc") %>%
  select(mean, min_n, mtry) %>%
  pivot_longer(min_n:mtry,
    values_to = "value",
    names_to = "parameter"
  ) %>%
  ggplot(aes(value, mean, color = parameter)) +
  geom_point(show.legend = FALSE) +
  facet_wrap(~parameter, scales = "free_x") +
  labs(x = NULL, y = "AUC")
```



```
autoplot(rf_tune, metric = "roc_auc")
```



```
rf_tune %>%
  show_best("roc_auc")


best_rf <- rf_tune %>%
  select_best("roc_auc")

final_rf_workflow <-
  rf_workflow %>%
  finalize_workflow(best_rf)

last_rf_fit <-
  final_rf_workflow %>%
  last_fit(df_split)

final_rf_fit <- extract_workflow(last_rf_fit)

rf_auc = validation(final_rf_fit, df_test)
```



```
# final_rf_fit %>%
#   extract_fit_parsnip() %>%
#   vip()
```

KNN

```
# knn_recipe <-
#   recipe(formula = sprintf("%s ~ .", outcome_column) %>% as.formula, data = df_train) %>%
#   step_string2factor(all_nominal_predictors()) %>%
#   step_novel(all_nominal_predictors()) %>%
#   step_zv(all_predictors()) %>%
#   step_dummy(all_nominal_predictors()) %>%
#   step_impute_mode(all_nominal_predictors()) %>%
#   step_impute_mean(all_numeric_predictors())
#
# knn_spec <-
#   nearest_neighbor(neighbors = tune(),
#                     weight_func = tune(),
#                     dist_power = tune()) %>%
#   set_mode("classification") %>%
#   set_engine("kknn")
#
# knn_grid <- grid_latin_hypercube(neighbors(),
#                                   weight_func(),
#                                   dist_power(),
#                                   size = 10)
#
# knn_workflow <-
#   workflow() %>%
#   add_recipe(knn_recipe) %>%
#   add_model(knn_spec)
#
```

```

# knn_tune <-
#   knn_workflow %>%
#   tune_grid(resamples = df_folds,
#             grid = knn_grid)
#
# knn_tune %>%
#   collect_metrics()
#
# autoplot(knn_tune, metric = "roc_auc")
#
# knn_tune %>%
#   show_best("roc_auc")
#
# best_knn <- knn_tune %>%
#   select_best("roc_auc")
#
# final_knn_workflow <-
#   knn_workflow %>%
#   finalize_workflow(best_knn)
#
# last_knn_fit <-
#   final_knn_workflow %>%
#   last_fit(df_split)
#
# final_knn_fit <- extract_workflow(last_knn_fit)
#
# knn_auc = validation(final_knn_fit, df_test)

```

SVM

```

svm_recipe <-
  recipe(formula = sprintf("%s ~ .", outcome_column) %>% as.formula, data = df_train) %>%
  step_string2factor(all_nominal_predictors()) %>%
  step_novel(all_nominal_predictors()) %>%
  step_zv(all_predictors()) %>%
  step_dummy(all_nominal_predictors()) %>%
  step_impute_mode(all_nominal_predictors()) %>%
  step_impute_mean(all_numeric_predictors())

svm_spec <-
  svm_rbf(cost = tune(), rbf_sigma = tune()) %>%
  set_mode("classification") %>%
  set_engine("kernlab")

svm_grid <- grid_latin_hypercube(cost(),
                                rbf_sigma(),
                                size = 10)

svm_workflow <-
  workflow() %>%
  add_recipe(svm_recipe) %>%
  add_model(svm_spec)

svm_tune <-
  svm_workflow %>%
  tune_grid(resamples = df_folds,
            grid = svm_grid)

svm_tune %>%


```

```
collect_metrics()
```

```
## # A tibble: 20 x 8
```

##	cost	rbf_sigma	.metric	.estimator	mean	n	std_err	.config
##	<dbl>	<dbl>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
## 1	0.00126	8.06e- 4	accuracy	binary	0.963	4	0.00309	Preprocessor1_Model101
## 2	0.00126	8.06e- 4	roc_auc	binary	0.636	4	0.0200	Preprocessor1_Model101
## 3	0.00455	1.91e- 7	accuracy	binary	0.963	4	0.00309	Preprocessor1_Model102
## 4	0.00455	1.91e- 7	roc_auc	binary	0.518	4	0.0271	Preprocessor1_Model102
## 5	0.0546	2.38e- 2	accuracy	binary	0.963	4	0.00309	Preprocessor1_Model103
## 6	0.0546	2.38e- 2	roc_auc	binary	0.634	4	0.00775	Preprocessor1_Model103
## 7	0.0733	1.48e- 6	accuracy	binary	0.963	4	0.00309	Preprocessor1_Model104
## 8	0.0733	1.48e- 6	roc_auc	binary	0.483	4	0.0282	Preprocessor1_Model104
## 9	9.75	1.85e-10	accuracy	binary	0.963	4	0.00309	Preprocessor1_Model105
## 10	9.75	1.85e-10	roc_auc	binary	0.511	4	0.0231	Preprocessor1_Model105
## 11	1.83	2.68e- 9	accuracy	binary	0.963	4	0.00309	Preprocessor1_Model106
## 12	1.83	2.68e- 9	roc_auc	binary	0.590	4	0.0182	Preprocessor1_Model106
## 13	21.1	2.54e- 5	accuracy	binary	0.961	4	0.00298	Preprocessor1_Model107
## 14	21.1	2.54e- 5	roc_auc	binary	0.484	4	0.0139	Preprocessor1_Model107
## 15	1.19	1.17e- 1	accuracy	binary	0.963	4	0.00309	Preprocessor1_Model108
## 16	1.19	1.17e- 1	roc_auc	binary	0.612	4	0.0111	Preprocessor1_Model108
## 17	0.0120	4.32e- 3	accuracy	binary	0.963	4	0.00309	Preprocessor1_Model109
## 18	0.0120	4.32e- 3	roc_auc	binary	0.611	4	0.0107	Preprocessor1_Model109
## 19	0.180	4.29e- 8	accuracy	binary	0.963	4	0.00309	Preprocessor1_Model110
## 20	0.180	4.29e- 8	roc_auc	binary	0.559	4	0.0116	Preprocessor1_Model110

```
autoplot(svm_tune, metric = "roc_auc")
```



```
svm_tune %>%
  show_best("roc_auc")
```

```
## # A tibble: 5 x 8
```

##	cost	rbf_sigma	.metric	.estimator	mean	n	std_err	.config
##	<dbl>	<dbl>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>


```
## 1 0.00126 0.000806      roc_auc binary      0.636      4 0.0200 Preprocessor1_Model01
## 2 0.0546  0.0238      roc_auc binary      0.634      4 0.00775 Preprocessor1_Model03
## 3 1.19    0.117      roc_auc binary      0.612      4 0.0111 Preprocessor1_Model08
## 4 0.0120  0.00432     roc_auc binary      0.611      4 0.0107 Preprocessor1_Model09
## 5 1.83    0.00000000268 roc_auc binary      0.590      4 0.0182 Preprocessor1_Model06
```


```
best_svm <- svm_tune %>%
  select_best("roc_auc")
```

```
final_svm_workflow <-
  svm_workflow %>%
  finalize_workflow(best_svm)
```

```
last_svm_fit <-
  final_svm_workflow %>%
  last_fit(df_split)
```

```
final_svm_fit <- extract_workflow(last_svm_fit)
```

```
svm_auc = validation(final_svm_fit, df_test)
```



```
## Confusion Matrix and Statistics
```

```
##
```

```
##
```

```
## test_predictions_class    0    1
##                0 4540  190
##                1    0    0
```

```
##
```

```
## Accuracy : 0.9598
```

```
## 95% CI : (0.9538, 0.9652)
```

```
## No Information Rate : 0.9598
```

```
## P-Value [Acc > NIR] : 0.5193
```

```
##
```


```
## Kappa : 0
```

```
##
## McNemar's Test P-Value : <2e-16
##
##      Sensitivity : 1.0000
##      Specificity : 0.0000
##      Pos Pred Value : 0.9598
##      Neg Pred Value :   NaN
##      Prevalence : 0.9598
##      Detection Rate : 0.9598
##      Detection Prevalence : 1.0000
##      Balanced Accuracy : 0.5000
##
##      'Positive' Class : 0
##
```

Models Comparison

```
df_auc <- tibble::tribble(
  ~Model, ~`AUC`, ~`Lower Limit`, ~`Upper Limit`,
  'XGBoost', as.numeric(xgboost_auc$auc), xgboost_auc$ci[1], xgboost_auc$ci[3],
  'GLM', as.numeric(glm_auc$auc), glm_auc$ci[1], glm_auc$ci[3],
  'Decision Tree', as.numeric(tree_auc$auc), tree_auc$ci[1], tree_auc$ci[3],
  'Random Forest', as.numeric(rf_auc$auc), rf_auc$ci[1], rf_auc$ci[3]
) %>%
  mutate(Target = outcome_column)

df_auc %>%
  ggplot(aes(x = Model, y = AUC, ymin = `Lower Limit`, ymax = `Upper Limit`)) +
    geom_point() +
    geom_errorbar()
```



```
saveRDS(df_auc, sprintf("../EDA/auxiliar/performance/%s_auc_result.RData", outcome_column))
```