

# Model Selection - death\_1year

Eduardo Yuki Yada

## Global parameters

```
k <- params$k # Number of folds for cross validation
grid_size <- params$grid_size # Number of parameter combination to tune on each model
repeats <- params$repeats
RUN_ALL_MODELS <- params$RUN_ALL_MODELS
Hmisc::list.tree(params)
```

```
## params = list 5 (952 bytes)
## . outcome_column = character 1= death_1year
## . k = double 1= 10
## . grid_size = double 1= 20
## . repeats = double 1= 2
## . RUN_ALL_MODELS = logical 1= TRUE
```

Minutes to run: 0

## Imports

```
library(tidyverse)
library(yaml)
library(tidymodels)
library(usemodels)
library(vip)
library(bonsai)
library(lightgbm)
library(caret)
library(pROC)

source("aux_functions.R")
predict <- stats::predict
```

Minutes to run: 0

## Loading data

```
load('dataset/processed_data.RData')
load('dataset/processed_dictionary.RData')

columns_list <- yaml.load_file("./auxiliar/columns_list.yaml")

outcome_column <- params$outcome_column
features_list <- params$features_list

df <- mutate(df, across(where(is.character), as.factor))
```

Minutes to run: 0.007

```
dir.create(file.path("./auxiliar/model_selection/hyperparameters/"),
          showWarnings = FALSE,
```

```

        recursive = TRUE)

dir.create(file.path("../auxiliar/model_selection/performance/"),
          showWarnings = FALSE,
          recursive = TRUE)

```

Minutes to run: 0

## Eligible features

```

cat_features_list = read_yaml(sprintf(
  "../auxiliar/significant_columns/categorical_%s.yaml",
  outcome_column
))

num_features_list = read_yaml(sprintf(
  "../auxiliar/significant_columns/numerical_%s.yaml",
  outcome_column
))

features_list = c(cat_features_list, num_features_list)

```

Minutes to run: 0

```

eligible_columns = df_names %>%
  filter(momento.aquisicao == 'Admissão t0') %>%
  .$variable.name

exception_columns = c('death_intraop', 'death_intraop_1', 'disch_outcomes_t0')

correlated_columns = c('year_procedure_1', # com year_adm_t0
  'age_surgery_1', # com age
  'admission_t0', # com admission_pre_t0_count
  'atb', # com meds_antimicrobianos
  'classe_meds_cardio_qtde', # com classe_meds_qtde
  'suporte_hemod', # com proced_invasivos_qtde,
  'radiografia', # com exames_imagem_qtde
  'ecg' # com metodos_graficos_qtde
)

eligible_features = eligible_columns %>%
  base::intersect(c(columns_list$categorical_columns, columns_list$numerical_columns)) %>%
  setdiff(c(exception_columns, correlated_columns))

features = base::intersect(eligible_features, features_list)

gluedown::md_order(features, seq = TRUE, pad = TRUE)

## 01. sex
## 02. age
## 03. education_level
## 04. underlying_heart_disease
## 05. heart_disease
## 06. nyha_basal
## 07. hypertension
## 08. prior_mi
## 09. heart_failure
## 10. af
## 11. cardiac_arrest
## 12. valvopathy
## 13. diabetes

```

## 14. renal\_failure  
## 15. hemodialysis  
## 16. stroke  
## 17. copd  
## 18. cancer  
## 19. comorbidities\_count  
## 20. procedure\_type\_1  
## 21. reop\_type\_1  
## 22. procedure\_type\_new  
## 23. cied\_final\_1  
## 24. cied\_final\_group\_1  
## 25. admission\_pre\_t0\_count  
## 26. admission\_pre\_t0\_180d  
## 27. year\_adm\_t0  
## 28. icu\_t0  
## 29. dialysis\_t0  
## 30. admission\_t0\_emergency  
## 31. aco  
## 32. antiarritmico  
## 33. ieca\_bra  
## 34. dva  
## 35. digoxina  
## 36. estatina  
## 37. diuretico  
## 38. vasodilatador  
## 39. insuf\_cardiaca  
## 40. espirolactona  
## 41. antiplaquetario\_ev  
## 42. insulina  
## 43. psicofarmacos  
## 44. antifungico  
## 45. antiviral  
## 46. classe\_meds\_qtde  
## 47. meds\_cardiovasc\_qtde  
## 48. meds\_antimicrobianos  
## 49. vni  
## 50. ventilacao\_mecanica  
## 51. transplante\_cardiaco  
## 52. cir\_toracica  
## 53. outros\_proced\_cirurgicos  
## 54. icp  
## 55. cateterismo  
## 56. cateter\_venoso\_central  
## 57. proced\_invasivos\_qtde  
## 58. transfusao  
## 59. interconsulta  
## 60. equipe\_multiprof  
## 61. holter  
## 62. teste\_esforco  
## 63. tilt\_teste  
## 64. metodos\_graficos\_qtde  
## 65. laboratorio  
## 66. cultura  
## 67. analises\_clinicas\_qtde  
## 68. citologia  
## 69. histopatologia\_qtde  
## 70. angio\_tc  
## 71. angiografia  
## 72. aortografia  
## 73. cintilografia  
## 74. ecocardiograma

```
## 75. endoscopia
## 76. flebografia
## 77. pet_ct
## 78. ultrassom
## 79. tomografia
## 80. ressonancia
## 81. exames_imagem_qtde
## 82. bic
## 83. hospital_stay
```

Minutes to run: 0

## Train test split (70%/30%)

```
set.seed(42)

if (outcome_column == 'readmission_30d') {
  df_split <- readRDS("./dataset/split_object.rds")
} else {
  df_split <- initial_split(df, prop = .7, strata = all_of(outcome_column))
}

df_train <- training(df_split) %>% dplyr::select(all_of(c(features, outcome_column)))
df_test <- testing(df_split) %>% dplyr::select(all_of(c(features, outcome_column)))

df_folds <- vfold_cv(df_train, v = k,
                     strata = all_of(outcome_column))
```

Minutes to run: 0.002

## Boosted Tree (XGBoost)

```
xgboost_recipe <-
  recipe(formula = sprintf("%s ~ .", outcome_column) %>% as.formula, data = df_train) %>%
  step_novel(all_nominal_predictors()) %>%
  step_unknown(all_nominal_predictors()) %>%
  step_other(all_nominal_predictors(), threshold = 0.05, other = ".merged") %>%
  step_dummy(all_nominal_predictors())

xgboost_spec <- boost_tree(
  trees = tune(),
  min_n = tune(),
  tree_depth = tune(),
  learn_rate = tune(),
) %>%
  set_engine("xgboost",
             nthread = 8) %>%
  set_mode("classification")

xgboost_grid <- grid_latin_hypercube(
  trees(range = c(25L, 150L)),
  min_n(range = c(2L, 100L)),
  tree_depth(range = c(2L, 15L)),
  learn_rate(range = c(-3, -1), trans = log10_trans()),
  size = grid_size
)

xgboost_workflow <-
  workflow() %>%
  add_recipe(xgboost_recipe) %>%
```

```

add_model(xgboost_spec)

xgboost_tune <-
  xgboost_workflow %>%
  tune_grid(resamples = df_folds,
            grid = xgboost_grid)


xgboost_tune %>%
  show_best("roc_auc")

## # A tibble: 5 x 10
##   trees min_n tree_depth learn_rate .metric .estimator mean      n std_err .config
##   <int> <int>    <int>    <dbl> <chr>   <chr>   <dbl> <int>  <dbl> <chr>
## 1    66     4         3    0.0689 roc_auc binary    0.800    10  0.0174 Preprocessor1_Model01
## 2   127    93         6    0.0872 roc_auc binary    0.793    10  0.0155 Preprocessor1_Model19
## 3    54    14        15    0.0555 roc_auc binary    0.789    10  0.0148 Preprocessor1_Model03
## 4    62    21        14    0.0482 roc_auc binary    0.789    10  0.0141 Preprocessor1_Model04
## 5   120     8        11    0.0203 roc_auc binary    0.787    10  0.0154 Preprocessor1_Model02

best_xgboost <- xgboost_tune %>%
  select_best("roc_auc")

autoplot(xgboost_tune, metric = "roc_auc")

```



```


final_xgboost_workflow <-
  xgboost_workflow %>%
  finalize_workflow(best_xgboost)

last_xgboost_fit <-
  final_xgboost_workflow %>%
  last_fit(df_split)


final_xgboost_fit <- extract_workflow(last_xgboost_fit)

```

```
xgboost_auc <- validation(final_xgboost_fit, df_test)
```



```
## [1] "Optimal Threshold: 0.04"
## Confusion Matrix and Statistics
##
##      reference
## data    0    1
##    0 3350   43
##    1 1219  118
##
##              Accuracy : 0.7332
##              95% CI   : (0.7203, 0.7458)
##    No Information Rate : 0.966
##    P-Value [Acc > NIR] : 1
##
##              Kappa   : 0.103
##
## Mcnemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.73320
##              Specificity : 0.73292
##              Pos Pred Value : 0.98733
##              Neg Pred Value : 0.08826
##              Prevalence   : 0.96596
##              Detection Rate : 0.70825
##              Detection Prevalence : 0.71734
##              Balanced Accuracy : 0.73306
##
##              'Positive' Class : 0
##
extract_vip(final_xgboost_fit, pred_wrapper = predict,
            reference_class = "0")
```



```
xgboost_parameters <- xgboost_tune %>%
  show_best("roc_auc", n = 1) %>%
  select(-.metric, -.estimator, -.config, -mean, -n, -std_err) %>%
  as.list
```

Minutes to run: 4.776

## Boosted Tree (LightGBM)

```
lightgbm_recipe <-
  recipe(formula = sprintf("%s ~ .", outcome_column) %>% as.formula, data = df_train) %>%
  step_novel(all_nominal_predictors()) %>%
  step_unknown(all_nominal_predictors()) %>%
  step_other(all_nominal_predictors(), threshold = 0.05, other = ".merged") %>%
  step_dummy(all_nominal_predictors())

lightgbm_spec <- boost_tree(
  trees = tune(),
  min_n = tune(),
  tree_depth = tune(),
  learn_rate = tune(),
  sample_size = 1
) %>%
  set_engine("lightgbm",
    nthread = 8) %>%
  set_mode("classification")

lightgbm_grid <- grid_latin_hypercube(
  trees(range = c(25L, 150L)),
  min_n(range = c(2L, 100L)),
  tree_depth(range = c(2L, 15L)),
  learn_rate(range = c(-3, -1), trans = log10_trans()),
  size = grid_size
```

```
)
```

```
lightgbm_workflow <-  
  workflow() %>%  
  add_recipe(lightgbm_recipe) %>%  
  add_model(lightgbm_spec)
```

```
lightgbm_tune <-  
  lightgbm_workflow %>%  
  tune_grid(resamples = df_folds,  
            grid = lightgbm_grid)
```


```
lightgbm_tune %>%  
  show_best("roc_auc")
```

```
## # A tibble: 5 x 10
```

```
##   trees min_n tree_depth learn_rate .metric .estimator  mean      n std_err .config  
##   <int> <int>    <int>    <dbl> <chr>   <chr>    <dbl> <int>  <dbl> <chr>  
## 1    99     3        6    0.0261 roc_auc binary   0.815    10 0.0122 Preprocessor1_Model01  
## 2    32    43       10    0.0652 roc_auc binary   0.809    10 0.0131 Preprocessor1_Model09  
## 3   105    73       13    0.0475 roc_auc binary   0.807    10 0.00846 Preprocessor1_Model15  
## 4   149    20        7    0.0123 roc_auc binary   0.806    10 0.0136 Preprocessor1_Model04  
## 5    53    23       11    0.0375 roc_auc binary   0.804    10 0.0136 Preprocessor1_Model05
```

```
best_lightgbm <- lightgbm_tune %>%  
  select_best("roc_auc")
```

```
autoplot(lightgbm_tune, metric = "roc_auc")
```



```
final_lightgbm_workflow <-  
  lightgbm_workflow %>%  
  finalize_workflow(best_lightgbm)
```

```
last_lightgbm_fit <-
```




```

final_lightgbm_workflow %>%
  last_fit(df_split)

final_lightgbm_fit <- extract_workflow(last_lightgbm_fit)

lightgbm_auc <- validation(final_lightgbm_fit, df_test)

```



```

## [1] "Optimal Threshold: 0.02"
## Confusion Matrix and Statistics
##
##      reference
## data    0    1
## 0 3318   40
## 1 1251  121
##
##              Accuracy : 0.7271
##              95% CI   : (0.7141, 0.7397)
##    No Information Rate : 0.966
##    P-Value [Acc > NIR] : 1
##
##              Kappa   : 0.1032
##
## Mcnemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.72620
##              Specificity : 0.75155
##              Pos Pred Value : 0.98809
##              Neg Pred Value : 0.08819
##              Prevalence   : 0.96596
##              Detection Rate : 0.70148
##              Detection Prevalence : 0.70994
##              Balanced Accuracy : 0.73888
##

```

```
##      'Positive' Class : 0
##
```

```
lightgbm_parameters <- lightgbm_tune %>%
  show_best("roc_auc", n = 1) %>%
  select(-.metric, -.estimator, -.config, -mean, -n, -std_err) %>%
  as.list
```


```
Hmisc::list.tree(lightgbm_parameters)
```

```
## lightgbm_parameters = list 4 (736 bytes)
## . trees = integer 1= 99
## . min_n = integer 1= 3
## . tree_depth = integer 1= 6
## . learn_rate = double 1= 0.026118
```

```
con <- file(sprintf('./auxiliar/model_selection/hyperparameters/%s.yaml', outcome_column), "w")
write_yaml(lightgbm_parameters, con)
close(con)
```

Minutes to run: 4.334

## Histogram of preds



0.013

Minutes to run:

## GLM

```
glmnet_recipe <-
  recipe(formula = sprintf("%s ~ .", outcome_column) %>% as.formula, data = df_train) %>%
  step_novel(all_nominal_predictors()) %>%
  step_unknown(all_nominal_predictors()) %>%
  step_other(all_nominal_predictors(), threshold = 0.05, other = ".merged") %>%
  step_dummy(all_nominal_predictors()) %>%
  step_zv(all_predictors()) %>%
  step_normalize(all_numeric_predictors())
```

```


glmnet_spec <-
  logistic_reg(penalty = 0) %>%
  set_mode("classification") %>%
  set_engine("glmnet")

glmnet_workflow <-
  workflow() %>%
  add_recipe(glmnet_recipe) %>%
  add_model(glmnet_spec)

glm_fit <- glmnet_workflow %>%
  fit(df_train)

glmnet_auc <- validation(glm_fit, df_test)

```



```

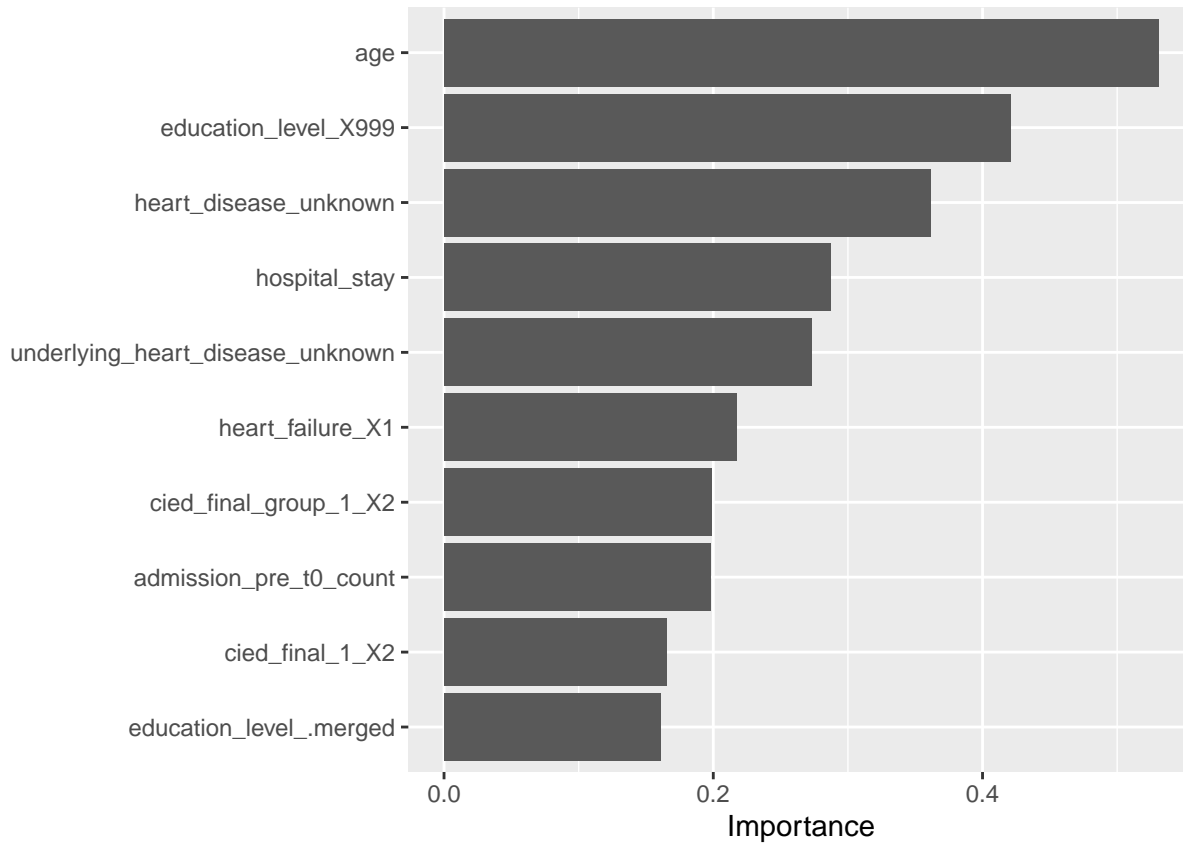
## [1] "Optimal Threshold: 0.03"
## Confusion Matrix and Statistics
##
##      reference
## data    0    1
## 0 3003   41
## 1 1566  120
##
##              Accuracy : 0.6603
##              95% CI   : (0.6466, 0.6738)
##    No Information Rate : 0.966
##    P-Value [Acc > NIR] : 1
##
##              Kappa   : 0.0723
##
##  Mcnemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.65726

```


```
##          Specificity : 0.74534
##          Pos Pred Value : 0.98653
##          Neg Pred Value : 0.07117
##          Prevalence : 0.96596
##          Detection Rate : 0.63488
##          Detection Prevalence : 0.64355
##          Balanced Accuracy : 0.70130
##
##          'Positive' Class : 0
##
```

```
pfun_glmnet <- function(object, newdata) predict(object, newx = newdata)

extract_vip(glm_fit, pred_wrapper = pfun_glmnet,
            reference_class = "1", method = 'model')
```



```
extract_vip(glm_fit, pred_wrapper = pfun_glmnet,
            reference_class = "1", method = 'permute')
```



Minutes to run:

1.923

## Decision Tree

```
tree_recipe <-
  recipe(formula = sprintf("%s ~ .", outcome_column) %>% as.formula, data = df_train) %>%
  step_novel(all_nominal_predictors()) %>%
  step_unknown(all_nominal_predictors()) %>%
  step_other(all_nominal_predictors(), threshold = 0.05, other = ".merged") %>%
  step_dummy(all_nominal_predictors()) %>%
  step_zv(all_predictors())

tree_spec <-
  decision_tree(cost_complexity = tune(),
                tree_depth = tune(),
                min_n = tune()) %>%
  set_mode("classification") %>%
  set_engine("rpart")

tree_grid <- grid_latin_hypercube(cost_complexity(),
                                   tree_depth(),
                                   min_n(),
                                   size = grid_size)


tree_workflow <-
  workflow() %>%
  add_recipe(tree_recipe) %>%
  add_model(tree_spec)

tree_tune <-
  tree_workflow %>%
  tune_grid(resamples = df_folds,
            grid = tree_grid)
```

```
tree_tune %>%
  collect_metrics()
```

```
## # A tibble: 40 x 9
##   cost_complexity tree_depth min_n .metric .estimator mean      n std_err .config
##         <dbl>      <int> <int> <chr>   <chr>    <dbl> <int>  <dbl> <chr>
## 1      1.49e-10        10     26 accuracy binary    0.964    10 0.00166 Preprocessor1_Model101
## 2      1.49e-10        10     26 roc_auc  binary    0.672    10 0.0147  Preprocessor1_Model101
## 3      4.39e-10        15     11 accuracy binary    0.956    10 0.00145 Preprocessor1_Model102
## 4      4.39e-10        15     11 roc_auc  binary    0.685    10 0.0169  Preprocessor1_Model102
## 5      4.48e- 5         11      6 accuracy binary    0.956    10 0.00145 Preprocessor1_Model103
## 6      4.48e- 5         11      6 roc_auc  binary    0.679    10 0.0250  Preprocessor1_Model103
## 7      3.67e- 7         13     19 accuracy binary    0.960    10 0.00162 Preprocessor1_Model104
## 8      3.67e- 7         13     19 roc_auc  binary    0.686    10 0.0168  Preprocessor1_Model104
## 9      9.28e- 6         10     30 accuracy binary    0.966    10 0.00147 Preprocessor1_Model105
## 10     9.28e- 6         10     30 roc_auc  binary    0.673    10 0.0145  Preprocessor1_Model105
## # ... with 30 more rows
## # i Use 'print(n = ...)' to see more rows
```

```
autoplot(tree_tune, metric = "roc_auc")
```



```
tree_tune %>%
  show_best("roc_auc")
```

```
## # A tibble: 5 x 9
##   cost_complexity tree_depth min_n .metric .estimator mean      n std_err .config
##         <dbl>      <int> <int> <chr>   <chr>    <dbl> <int>  <dbl> <chr>
## 1      0.000111         9      8 roc_auc binary    0.721    10 0.0166 Preprocessor1_Model113
## 2      0.00116         8      6 roc_auc binary    0.699    10 0.0167 Preprocessor1_Model110
## 3      0.000488        13     13 roc_auc binary    0.696    10 0.0196 Preprocessor1_Model106
## 4      0.000000752      4      2 roc_auc binary    0.695    10 0.0140 Preprocessor1_Model115
## 5      0.000000367     13     19 roc_auc binary    0.686    10 0.0168 Preprocessor1_Model104
```

```

best_tree <- tree_tune %>%
  select_best("roc_auc")


final_tree_workflow <-
  tree_workflow %>%
  finalize_workflow(best_tree)

last_tree_fit <-
  final_tree_workflow %>%
  last_fit(df_split)

final_tree_fit <- extract_workflow(last_tree_fit)

tree_auc <- validation(final_tree_fit, df_test)

```




```

## [1] "Optimal Threshold: 0.03"
## Confusion Matrix and Statistics
##
##      reference
## data    0    1
## 0 3525   68
## 1 1044   93
##
##              Accuracy : 0.7649
##              95% CI : (0.7526, 0.7769)
##    No Information Rate : 0.966
##    P-Value [Acc > NIR] : 1
##
##              Kappa : 0.089
##
##  McNemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.77150

```

```
##          Specificity : 0.57764
##          Pos Pred Value : 0.98107
##          Neg Pred Value : 0.08179
##          Prevalence : 0.96596
##          Detection Rate : 0.74524
##          Detection Prevalence : 0.75962
##          Balanced Accuracy : 0.67457
##
##          'Positive' Class : 0
##
```

```
extract_vip(final_tree_fit, pred_wrapper = predict,
            reference_class = "0", use_matrix = FALSE,
            method = 'model')
```



```
# extract_vip(final_tree_fit, pred_wrapper = predict,
#             reference_class = "1", use_matrix = FALSE,
#             method = 'permute')
```

Minutes to run: 4.538

## Random Forest

```
rf_recipe <-
  recipe(formula = sprintf("%s ~ .", outcome_column) %>% as.formula,
        data = df_train) %>%
  step_novel(all_nominal_predictors()) %>%
  step_unknown(all_nominal_predictors()) %>%
  step_other(all_nominal_predictors(), threshold = 0.05, other = ".merged") %>%
  step_dummy(all_nominal_predictors()) %>%
  step_zv(all_predictors()) %>%
  step_impute_mean(all_numeric_predictors())

rf_spec <-
```



```

rand_forest(mtry = tune(),
            trees = tune(),
            min_n = tune()) %>%
set_mode("classification") %>%
set_engine("randomForest",
            probability = TRUE,
            nthread = 8)

rf_grid <- grid_latin_hypercube(mtry(range = c(1L, 50L)),
                               trees(range = c(100L, 300L)),
                               min_n(),
                               size = grid_size)

rf_workflow <-
  workflow() %>%
  add_recipe(rf_recipe) %>%
  add_model(rf_spec)


rf_tune <-
  rf_workflow %>%
  tune_grid(resamples = df_folds,
            grid = rf_grid)

rf_tune %>%
  collect_metrics()

## # A tibble: 40 x 9
##   mtry trees min_n .metric .estimator mean n std_err .config
##   <int> <int> <int> <chr>   <chr>   <dbl> <int>  <dbl> <chr>
## 1    49   144    39 accuracy binary   0.968    10 0.00141 Preprocessor1_Model01
## 2    49   144    39 roc_auc  binary   0.734    10 0.0136  Preprocessor1_Model01
## 3    18   277    33 accuracy binary   0.968    10 0.00143 Preprocessor1_Model02
## 4    18   277    33 roc_auc  binary   0.742    10 0.0181  Preprocessor1_Model02
## 5    46   224     7 accuracy binary   0.967    10 0.00146 Preprocessor1_Model03
## 6    46   224     7 roc_auc  binary   0.761    10 0.0134  Preprocessor1_Model03
## 7    34   216    24 accuracy binary   0.968    10 0.00143 Preprocessor1_Model04
## 8    34   216    24 roc_auc  binary   0.762    10 0.0148  Preprocessor1_Model04
## 9    43   106    12 accuracy binary   0.967    10 0.00139 Preprocessor1_Model05
## 10   43   106    12 roc_auc  binary   0.751    10 0.0131  Preprocessor1_Model05
## # ... with 30 more rows
## # i Use 'print(n = ...)' to see more rows

autoplot(rf_tune, metric = "roc_auc")

```



```
rf_tune %>%
  show_best("roc_auc")
```

```
## # A tibble: 5 x 9
##   mtry trees min_n .metric .estimator mean     n std_err .config
##   <int> <int> <int> <chr>   <chr>   <dbl> <int>   <dbl> <chr>
## 1    14   188     6 roc_auc binary  0.775    10  0.0123 Preprocessor1_Model15
## 2    18   252    16 roc_auc binary  0.769    10  0.0151 Preprocessor1_Model108
## 3    32   267    11 roc_auc binary  0.768    10  0.0147 Preprocessor1_Model113
## 4    24   249    18 roc_auc binary  0.767    10  0.0128 Preprocessor1_Model118
## 5    39   233     9 roc_auc binary  0.766    10  0.0141 Preprocessor1_Model119
```


```
best_rf <- rf_tune %>%
  select_best("roc_auc")
```

```
final_rf_workflow <-
  rf_workflow %>%
  finalize_workflow(best_rf)
```

```
last_rf_fit <-
  final_rf_workflow %>%
  last_fit(df_split)
```

```
final_rf_fit <- extract_workflow(last_rf_fit)
```


```
rf_auc <- validation(final_rf_fit, df_test)
```



```
## [1] "Optimal Threshold: 0.02"
## Confusion Matrix and Statistics
##
##      reference
## data    0    1
##    0 2932   32
##    1 1637  129
##
##              Accuracy : 0.6471
##              95% CI   : (0.6333, 0.6608)
##    No Information Rate : 0.966
##    P-Value [Acc > NIR] : 1
##
##              Kappa : 0.0763
##
##  McNemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.64172
##              Specificity : 0.80124
##    Pos Pred Value : 0.98920
##    Neg Pred Value : 0.07305
##    Prevalence : 0.96596
##    Detection Rate : 0.61987
##    Detection Prevalence : 0.62664
##    Balanced Accuracy : 0.72148
##
##    'Positive' Class : 0
##
```

```
pfun_rf <- function(object, newdata) predict(object, data = newdata)

extract_vip(final_rf_fit, pred_wrapper = predict,
             reference_class = "1", use_matrix = FALSE,
             method = 'model')
```




```
# extract_vip(final_rf_fit, pred_wrapper = predict,
#             reference_class = "1", use_matrix = FALSE,
#             method = 'permute')
```

Minutes to run: 62.738

## Models Comparison

```
if (RUN_ALL_MODELS) {
  df_auc <- tibble::tribble(
    ~Model, ~`AUC`, ~`Lower Limit`, ~`Upper Limit`,
    'XGBoost', as.numeric(xgboost_auc$auc), xgboost_auc$ci[1], xgboost_auc$ci[3],
    'LightGBM', as.numeric(lightgbm_auc$auc), lightgbm_auc$ci[1], lightgbm_auc$ci[3],
    'GLM', as.numeric(glmnet_auc$auc), glmnet_auc$ci[1], glmnet_auc$ci[3],
    'Decision Tree', as.numeric(tree_auc$auc), tree_auc$ci[1], tree_auc$ci[3],
    'Random Forest', as.numeric(rf_auc$auc), rf_auc$ci[1], rf_auc$ci[3]
  ) %>%
    mutate(Target = outcome_column)
} else {
  df_auc <- tibble::tribble(
    ~Model, ~`AUC`, ~`Lower Limit`, ~`Upper Limit`,
    'LightGBM', as.numeric(lightgbm_auc$auc), lightgbm_auc$ci[1], lightgbm_auc$ci[3]
  ) %>%
    mutate(Target = outcome_column)
}

df_auc %>%
  ggplot(aes(x = Model, y = AUC, ymin = `Lower Limit`, ymax = `Upper Limit`)) +
    geom_point() +
    geom_errorbar()
```



```
write_csv(df_auc, sprintf("./auxiliar/model_selection/performance/%s.csv", outcome_column))
```

Minutes to run: 0.002