

# Final Model - death\_180days

Eduardo Yuki Yada

## Global parameters

```
k <- 5 # Number of folds for cross validation
grid_size <- 30 # Number of parameter combination to tune on each model
max_auc_loss <- 0.01 # Max accepted loss of AUC for reducing num of features
```

## Imports

```
library(tidyverse)
library(yaml)
library(tidymodels)
library(usemodels)
library(vip)
library(kableExtra)
library(SHAPforxgboost)
library(xgboost)
library(Matrix)
library(mltools)
library(bonsai)
library(lightgbm)
library(pROC)
library(caret)
library(themis)

source("aux_functions.R")

select <- dplyr::select
```

## Loading data

```
load('dataset/processed_data.RData')
load('dataset/processed_dictionary.RData')

columns_list <- yaml.load_file("./auxiliar/columns_list.yaml")

outcome_column <- params$outcome_column
features_list <- params$features_list

df[columns_list$outcome_columns] <- lapply(df[columns_list$outcome_columns], factor)
df <- mutate(df, across(where(is.character), as.factor))

dir.create(file.path("./auxiliar/final_model/hyperparameters/"),
          showWarnings = FALSE,
          recursive = TRUE)

dir.create(file.path("./auxiliar/final_model/performance/"),
          showWarnings = FALSE,
          recursive = TRUE)
```

```

dir.create(file.path("./auxiliar/final_model/selected_features/"),
  showWarnings = FALSE,
  recursive = TRUE)

```

## Eligible features

```

cat_features_list = readRDS(sprintf(
  "./auxiliar/significant_columns/categorical_%s.rds",
  outcome_column
))

num_features_list = readRDS(sprintf(
  "./auxiliar/significant_columns/numerical_%s.rds",
  outcome_column
))

features_list = c(cat_features_list, num_features_list)

eligible_columns <- df_names %>%
  filter(momento.aquisicao == 'Admissão t0') %>%
  .$variable.name

exception_columns <- c('death_intraop', 'death_intraop_1', 'disch_outcomes_t0')

correlated_columns = c('year_procedure_1', # com year_adm_t0
  'age_surgery_1', # com age
  'admission_t0', # com admission_pre_t0_count
  'atb', # com meds_antimicrobianos
  'classe_meds_cardio_qtde', # com classe_meds_qtde
  'suporte_hemod', # com proced_invasivos_qtde,
  'radiografia', # com exames_imagem_qtde
  'ecg' # com metodos_graficos_qtde
  )

eligible_features <- eligible_columns %>%
  base::intersect(c(columns_list$categorical_columns, columns_list$numerical_columns)) %>%
  setdiff(c(exception_columns, correlated_columns))

if (is.null(features_list)) {
  features = eligible_features
} else {
  features = base::intersect(eligible_features, features_list)
}

```

Starting features:

```
gluedown::md_order(features, seq = TRUE, pad = TRUE)
```

1. sex
2. age
3. education\_level
4. underlying\_heart\_disease
5. heart\_disease
6. nyha\_basal
7. hypertension
8. prior\_mi
9. heart\_failure
10. af
11. cardiac\_arrest
12. valvopathy

13. diabetes  
14. renal\_failure  
15. hemodialysis  
16. stroke  
17. copd  
18. cancer  
19. comorbidities\_count  
20. procedure\_type\_1  
21. reop\_type\_1  
22. procedure\_type\_new  
23. cied\_final\_1  
24. cied\_final\_group\_1  
25. admission\_pre\_t0\_count  
26. admission\_pre\_t0\_180d  
27. year\_adm\_t0  
28. icu\_t0  
29. dialysis\_t0  
30. admission\_t0\_emergency  
31. aco  
32. antiarritmico  
33. ieca\_bra  
34. dva  
35. digoxina  
36. estatina  
37. diuretico  
38. vasodilatador  
39. insuf\_cardiaca  
40. espironolactona  
41. antiplaquetario\_ev  
42. insulina  
43. psicofarmacos  
44. antifungico  
45. classe\_meds\_qtde  
46. meds\_cardiovasc\_qtde  
47. meds\_antimicrobianos  
48. vni  
49. ventilacao\_mecanica  
50. transplante\_cardiaco  
51. outros\_proced\_cirurgicos  
52. icp  
53. cateterismo  
54. cateter\_venoso\_central  
55. proced\_invasivos\_qtde  
56. transfusao  
57. interconsulta  
58. equipe\_multiprof  
59. holter  
60. teste\_esforco  
61. metodos\_graficos\_qtde  
62. laboratorio  
63. cultura  
64. analises\_clinicas\_qtde  
65. citologia  
66. histopatologia\_qtde  
67. angiografia  
68. aortografia  
69. arteriografia  
70. cintilografia  
71. ecocardiograma  
72. endoscopia  
73. ultrassom

74. tomografia  
 75. ressonancia  
 76. exames\_imagem\_qtde  
 77. bic  
 78. hospital\_stay

## Train test split (70%/30%)

```

set.seed(42)

if (outcome_column == 'readmission_30d') {
  df_split <- readRDS("dataset/split_object.rds")
} else {
  df_split <- initial_split(df, prop = .7, strata = all_of(outcome_column))
}

df_train <- training(df_split) %>% select(all_of(c(features, outcome_column)))
df_test <- testing(df_split) %>% select(all_of(c(features, outcome_column)))

df_folds <- vfold_cv(df_train, v = k,
                      strata = all_of(outcome_column))

```

## Feature Selection

```

model_fit_wf <- function(df_train, features, outcome_column, hyperparameters){
  model_recipe <-
    recipe(formula = sprintf("%s ~ .", outcome_column) %>% as.formula,
           data = df_train %>% select(all_of(c(features, outcome_column)))) %>%
    step_novel(all_nominal_predictors()) %>%
    step_unknown(all_nominal_predictors()) %>%
    step_other(all_nominal_predictors(), threshold = 0.05, other = ".merged")

  model_spec <-
    do.call(boost_tree, hyperparameters) %>%
    set_engine("lightgbm") %>%
    set_mode("classification")

  model_workflow <-
    workflow() %>%
    add_recipe(model_recipe) %>%
    add_model(model_spec)

  model_fit_rs <- model_workflow %>%
    fit_resamples(df_folds)

  model_fit <- model_workflow %>%
    fit(df_train)

  model_auc <- validation(model_fit, df_test, plot = F)

  raw_model <- parsnip::extract_fit_engine(model_fit)

  feature_importance <- lgb.importance(raw_model, percentage = TRUE)

  cv_results <- collect_metrics(model_fit_rs) %>% filter(.metric == 'roc_auc')

  return(
    list(
      cv_auc = cv_results$mean,

```

```

    cv_auc_std_err = cv_results$std_err,
    importance = feature_importance,
    auc = as.numeric(model_auc$auc),
    auc_lower = model_auc$ci[1],
    auc_upper = model_auc$ci[3]
  )
)
}

hyperparameters <- readRDS(
  sprintf(
    "./auxiliar/model_selection/hyperparameters/lightgbm_%s.rds",
    outcome_column
  )
)

hyperparameters$sample_size <- 1

full_model <- model_fit_wf(df_train, features, outcome_column, hyperparameters)

sprintf('Full Model CV Train AUC: %.3f' ,full_model$cv_auc)

## [1] "Full Model CV Train AUC: 0.778"
sprintf('Full Model Test AUC: %.3f' ,full_model$auc)

## [1] "Full Model Test AUC: 0.808"

Features with zero importance on the initial model:

unimportant_features <- setdiff(features, full_model$importance$Feature)

unimportant_features %>%
  gluedown::md_order()

1. education_level
2. underlying_heart_disease
3. heart_disease
4. hemodialysis
5. cied_final_1
6. dialysis_t0
7. vni
8. transplante_cardiaco
9. outros_proced_cirurgicos
10. icp
11. cateter_venoso_central
12. transfusao
13. teste_esforco
14. angiografia
15. aortografia
16. arteriografia

trimmed_features <- full_model$importance$Feature
hyperparameters$mtry <- min(hyperparameters$mtry, length(trimmed_features))
trimmed_model <- model_fit_wf(df_train, trimmed_features,
                                outcome_column, hyperparameters)

sprintf('Trimmed Model CV Train AUC: %.3f' ,trimmed_model$cv_auc)

## [1] "Trimmed Model CV Train AUC: 0.776"
sprintf('Trimmed Model Test AUC: %.3f' ,trimmed_model$auc)

## [1] "Trimmed Model Test AUC: 0.803"

```

```

selection_results <- tibble::tribble(
  ~`Tested Feature`, ~`Dropped`, ~`Number of Features`, ~`CV AUC`, ~`CV AUC Std Error`, ~`Total AUC Loss`, ~`Instant AUC Loss`,
  'None', TRUE, length(features), full_model$cv_auc, full_model$cv_auc_std_err, 0, 0
)

whitelist <- c()

if (full_model$cv_auc - trimmed_model$cv_auc < max_auc_loss) {
  current_features <- trimmed_features
  current_model <- trimmed_model
  current_auc_loss <- full_model$cv_auc - current_model$cv_auc
  instant_auc_loss <- full_model$cv_auc - current_model$cv_auc

  selection_results <- selection_results %>%
    add_row(`Tested Feature` = 'All unimportant',
            `Dropped` = TRUE,
            `Number of Features` = length(trimmed_features),
            `CV AUC` = current_model$cv_auc,
            `CV AUC Std Error` = current_model$cv_auc_std_err,
            `Total AUC Loss` = current_auc_loss,
            `Instant AUC Loss` = instant_auc_loss)
} else {
  current_features <- features
  current_model <- full_model
  current_auc_loss <- 0
}

while (current_auc_loss < max_auc_loss | mean(current_features %in% whitelist) == 1) {
  current_least_important <-
    tail(setdiff(current_model$importance$Feature, whitelist), 1)
  test_features <-
    setdiff(current_features, current_least_important)
  hyperparameters$mtry <-
    min(hyperparameters$mtry, length(test_features))
  current_model <-
    model_fit_wf(df_train, test_features, outcome_column, hyperparameters)
  instant_auc_loss <-
    tail(selection_results %>% filter(Dropped) %>% .[["CV AUC"]], n = 1) - current_model$cv_auc
  current_auc_loss <- full_model$cv_auc - current_model$cv_auc

  if (instant_auc_loss < max_auc_loss / 5 &
      current_auc_loss < max_auc_loss) {
    dropped <- TRUE
    current_features <- test_features
  } else {
    dropped <- FALSE
    whitelist <- c(whitelist, current_least_important)
  }

  selection_results <- selection_results %>%
    add_row(
      `Tested Feature` = current_least_important,
      `Dropped` = dropped,
      `Number of Features` = length(test_features),
      `CV AUC` = current_model$cv_auc,
      `CV AUC Std Error` = current_model$cv_auc_std_err,
      `Total AUC Loss` = current_auc_loss,
      `Instant AUC Loss` = instant_auc_loss
    )
}

print(c(

```

```

length(current_features),
round(current_auc_loss, 4),
round(instant_auc_loss, 4),
current_least_important
))
}

## [1] "61"           "0.0026"        "0.001"          "antiplaquetario_ev"
## [1] "61"           "0.0085"        "0.0058"         "procedure_type_1"
## [1] "60"           "0.0026"        "0"              "bic"
## [1] "59"           "0.0027"        "1e-04"          "cardiac_arrest"
## [1] "59"           "0.0055"        "0.0028"         "nyha_basal"
## [1] "59"           "0.005"         "0.0022"         "insulina"
## [1] "58"           "0.0024"        "-3e-04"         "aco"
## [1] "57"           "0.0033"        "9e-04"          "valvopathy"
## [1] "57"           "0.0092"        "0.0059"         "antifungico"
## [1] "56"           "0.0047"        "0.0014"         "holter"
## [1] "55"           "0.0059"        "0.0012"         "histopatologia_qtde"
## [1] "54"           "0.0064"        "4e-04"          "ventilacao_mecanica"
## [1] "53"           "0.0078"        "0.0015"         "reop_type_1"
## [1] "52"           "0.0096"        "0.0017"         "hypertension"
## [1] "51"           "0.0071"        "-0.0024"        "diabetes"
## [1] "51"           "0.0116"        "0.0044"         "ressonancia"

selection_results %>%
  rename(Features = `Number of Features`) %>%
  niceFormatting(digits = 4, label = 1)

```

Table 1:

Tested Feature	Dropped	Features	CV AUC	CV AUC Std Error	Total AUC Loss	Instant AUC Loss
None	TRUE	78	0.7778	0.0186	0.0000	0.0000
All unimportant	TRUE	62	0.7762	0.0170	0.0016	0.0016
antiplaquetario_ev	TRUE	61	0.7751	0.0168	0.0026	0.0010
procedure_type_1	FALSE	60	0.7693	0.0169	0.0085	0.0058
bic	TRUE	60	0.7751	0.0186	0.0026	0.0000
cardiac_arrest	TRUE	59	0.7750	0.0180	0.0027	0.0001
nyha_basal	FALSE	58	0.7722	0.0175	0.0055	0.0028
insulina	FALSE	58	0.7728	0.0154	0.0050	0.0022
aco	TRUE	58	0.7753	0.0173	0.0024	-0.0003
valvopathy	TRUE	57	0.7745	0.0170	0.0033	0.0009
antifungico	FALSE	56	0.7685	0.0167	0.0092	0.0059
holter	TRUE	56	0.7731	0.0188	0.0047	0.0014
histopatologia_qtde	TRUE	55	0.7718	0.0186	0.0059	0.0012
ventilacao_mecanica	TRUE	54	0.7714	0.0167	0.0064	0.0004
reop_type_1	TRUE	53	0.7699	0.0159	0.0078	0.0015
hypertension	TRUE	52	0.7682	0.0155	0.0096	0.0017
diabetes	TRUE	51	0.7707	0.0156	0.0071	-0.0024
ressonancia	FALSE	50	0.7662	0.0156	0.0116	0.0044

```

if (exists('last_feature_dropped')) {
  selected_features <- c(current_features, last_feature_dropped)
} else {
  selected_features <- current_features
}

con <- file(sprintf('./auxiliar/final_model/selected_features/%s.yaml', outcome_column), "w")
write_yaml(selected_features, con)
close(con)

```

```

feature_selected_model <- model_fit_wf(df_train, selected_features,
                                         outcome_column, hyperparameters)

sprintf('Selected Model CV Train AUC: %.3f', feature_selected_model$cv_auc)

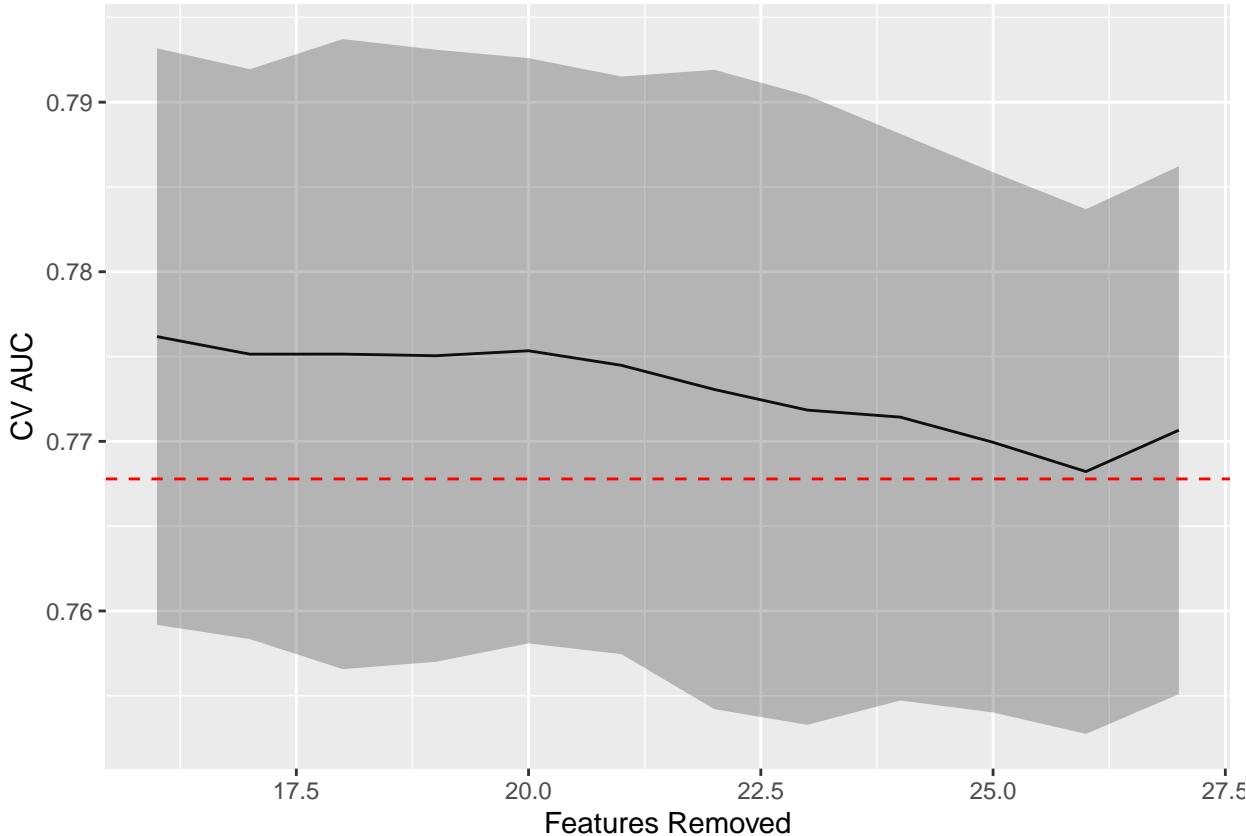
## [1] "Selected Model CV Train AUC: 0.768"
sprintf('Selected Model Test AUC: %.3f', feature_selected_model$auc)

## [1] "Selected Model Test AUC: 0.782"

selection_results <- selection_results %>%
  filter(`Number of Features` < length(features)) %>%
  mutate(`Features Removed` = length(features) - `Number of Features`,
        `CV AUC Low` = `CV AUC` - `CV AUC Std Error`,
        `CV AUC High` = `CV AUC` + `CV AUC Std Error`)

selection_results %>%
  filter(Dropped) %>%
  ggplot(aes(x = `Features Removed`, y = `CV AUC`,
             ymin = `CV AUC Low`, ymax = `CV AUC High`)) +
  geom_line() +
  geom_ribbon(alpha = .3) +
  geom_hline(yintercept = full_model$cv_auc - max_auc_loss,
             linetype = "dashed", color = "red")

```



## Hyperparameter tuning

Selected features:

```
gluedown::md_order(selected_features, seq = TRUE, pad = TRUE)
```

1. hospital\_stay
2. age

3. laboratorio  
 4. vasodilatador  
 5. espironolactona  
 6. year\_adm\_t0  
 7. ieca\_bra  
 8. analises\_clinicas\_qtde  
 9. admission\_pre\_t0\_count  
 10. comorbidities\_count  
 11. icu\_t0  
 12. equipe\_multiprof  
 13. meds\_cardiovasc\_qtde  
 14. meds\_antimicrobianos  
 15. diuretico  
 16. classe\_meds\_qtde  
 17. antiaritmico  
 18. estatina  
 19. admission\_pre\_t0\_180d  
 20. dva  
 21. insuf\_cardiaca  
 22. psicofarmacos  
 23. exames\_imagem\_qtde  
 24. ecocardiograma  
 25. stroke  
 26. metodos\_graficos\_qtde  
 27. cied\_final\_group\_1  
 28. admission\_t0\_emergency  
 29. cateterismo  
 30. citologia  
 31. interconsulta  
 32. proced\_invasivos\_qtde  
 33. cintilografia  
 34. renal\_failure  
 35. ultrassom  
 36. af  
 37. sex  
 38. cancer  
 39. copd  
 40. tomografia  
 41. cultura  
 42. prior\_mi  
 43. ressonancia  
 44. insulina  
 45. heart\_failure  
 46. procedure\_type\_new  
 47. digoxina  
 48. antifungico  
 49. procedure\_type\_1  
 50. nyha Basal  
 51. endoscopia

## Standard

```

lightgbm_recipe <-
  recipe(formula = sprintf("%s ~ .", outcome_column) %>% as.formula,
         data = df_train %>% select(all_of(c(selected_features, outcome_column)))) %>%
  step_novel(all_nominal_predictors()) %>%
  step_unknown(all_nominal_predictors()) %>%
  step_other(all_nominal_predictors(), threshold = 0.05, other = ".merged") %>%
  step_dummy(all_nominal_predictors())

```

```

lightgbm_smote_recipe <-

```

```

recipe(formula = sprintf("%s ~ .", outcome_column) %>% as.formula,
       data = df_train %>% select(all_of(c(selected_features, outcome_column)))) %>%
step_novel(all_nominal_predictors()) %>%
step_unknown(all_nominal_predictors()) %>%
step_other(all_nominal_predictors(), threshold = 0.05, other = ".merged") %>%
step_dummy(all_nominal_predictors()) %>%
step_impute_mean(all_numeric_predictors()) %>%
step_smote(!sym(outcome_column))

lightgbm_upsample_recipe <-
  recipe(formula = sprintf("%s ~ .", outcome_column) %>% as.formula,
         data = df_train %>% select(all_of(c(selected_features, outcome_column)))) %>%
  step_novel(all_nominal_predictors()) %>%
  step_unknown(all_nominal_predictors()) %>%
  step_other(all_nominal_predictors(), threshold = 0.05, other = ".merged") %>%
  step_dummy(all_nominal_predictors()) %>%
  step_upsample(!sym(outcome_column))

lightgbm_tuning <- function(recipe) {

  lightgbm_spec <- boost_tree(
    mtry = tune(),
    trees = tune(),
    min_n = tune(),
    tree_depth = tune(),
    learn_rate = tune(),
    loss_reduction = tune(),
    sample_size = 1.0
  ) %>%
    set_engine("lightgbm") %>%
    set_mode("classification")

  lightgbm_grid <- grid_latin_hypercube(
    mtry(range = c(1L, length(selected_features))),
    trees(range = c(100L, 300L)),
    min_n(),
    tree_depth(),
    learn_rate(),
    loss_reduction(),
    size = grid_size
  )

  lightgbm_workflow <-
    workflow() %>%
    add_recipe(recipe) %>%
    add_model(lightgbm_spec)

  lightgbm_tune <-
    lightgbm_workflow %>%
    tune_grid(resamples = df_folds,
              grid = lightgbm_grid)

  lightgbm_tune %>%
    show_best("roc_auc") %>%
    niceFormatting(digits = 5, label = 4)

  best_lightgbm <- lightgbm_tune %>%
    select_best("roc_auc")

  lightgbm_tune %>%
    collect_metrics() %>%

```

```

filter(.metric == "roc_auc") %>%
select(mean, mtry:tree_depth) %>%
pivot_longer(mtry:tree_depth,
             values_to = "value",
             names_to = "parameter"
) %>%
ggplot(aes(value, mean, color = parameter)) +
geom_point(alpha = 0.8, show.legend = FALSE) +
facet_wrap(~parameter, scales = "free_x") +
labs(x = NULL, y = "AUC")

final_lightgbm_workflow <-
lightgbm_workflow %>%
finalize_workflow(best_lightgbm)

last_lightgbm_fit <-
final_lightgbm_workflow %>%
last_fit(df_split)

final_lightgbm_fit <- extract_workflow(last_lightgbm_fit)

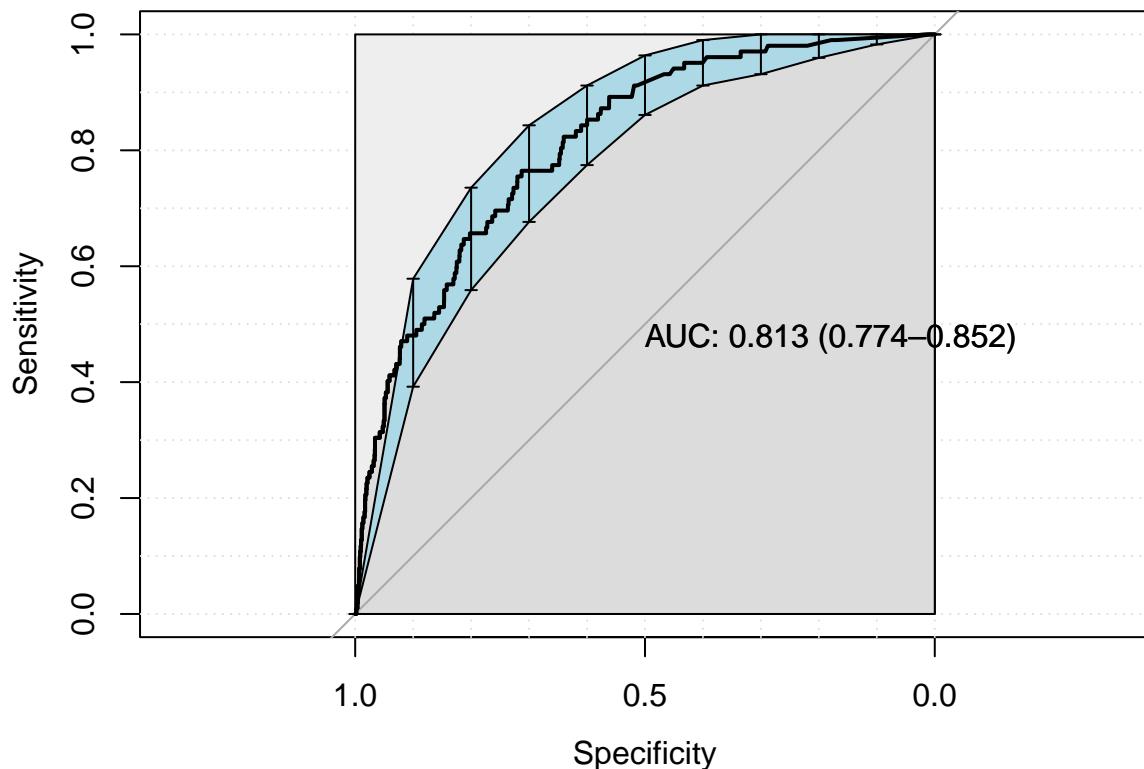
lightgbm_auc <- validation(final_lightgbm_fit, df_test)

lightgbm_parameters <- lightgbm_tune %>%
show_best("roc_auc", n = 1) %>%
select(trees, mtry, min_n, tree_depth, learn_rate, loss_reduction) %>%
as.list

return(list(auc = as.numeric(lightgbm_auc$auc),
            auc_lower = lightgbm_auc$ci[1],
            auc_upper = lightgbm_auc$ci[3],
            parameters = lightgbm_parameters,
            fit = final_lightgbm_fit))
}

standard_results <- lightgbm_tuning(lightgbm_recipe)

```



```

## [1] "Optimal Threshold: 0.02"
## Confusion Matrix and Statistics
##
##      reference
## data      0      1
##   0 3299    24
##   1 1329    78
##
##                  Accuracy : 0.714
##                  95% CI : (0.7008, 0.7268)
##      No Information Rate : 0.9784
##      P-Value [Acc > NIR] : 1
##
##                  Kappa : 0.0658
##
## Mcnemar's Test P-Value : <2e-16
##
##      Sensitivity : 0.71283
##      Specificity : 0.76471
##      Pos Pred Value : 0.99278
##      Neg Pred Value : 0.05544
##      Prevalence : 0.97844
##      Detection Rate : 0.69746
##      Detection Prevalence : 0.70254
##      Balanced Accuracy : 0.73877
##
##      'Positive' Class : 0
##
# smote_results <- lightgbm_tuning(lightgbm_smote_recipe)
# upsample_results <- lightgbm_tuning(lightgbm_upsample_recipe)

final_lightgbm_fit <- standard_results$fit
lightgbm_parameters <- standard_results$parameters

```

```

# saveRDS(
#   lightgbm_parameters,
#   file = sprintf(
#     "./auxiliar/final_model/hyperparameters/lightgbm_%s.rds",
#     outcome_column
#   )
# )

```

## SHAP values

```

lightgbm_model <- parsnip::extract_fit_engine(final_lightgbm_fit)

trained_rec <- prep(lightgbm_recipe, training = df_train)

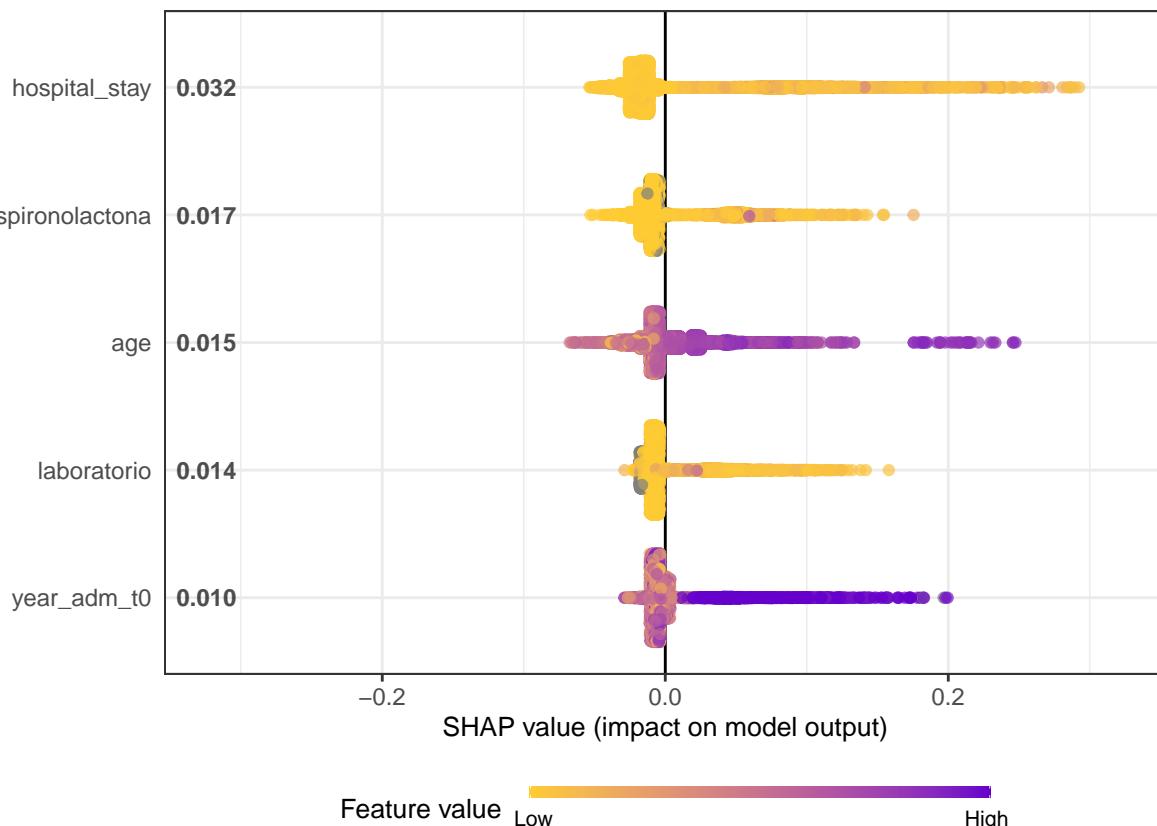
df_train_baked <- bake(trained_rec, new_data = df_train)
df_test_baked <- bake(trained_rec, new_data = df_test)

matrix_train <- as.matrix(df_train_baked %>% select(-all_of(outcome_column)))
matrix_test <- as.matrix(df_test_baked %>% select(-all_of(outcome_column)))

n_plots <- min(5, length(selected_features))

shap.plot.summary.wrap1(model = lightgbm_model, X = matrix_train,
                       top_n = n_plots, dilute = F)

```



```

shap <- shap.prep(lightgbm_model, X_train = matrix_test)

for (x in shap.importance(shap, names_only = TRUE)[1:n_plots]) {
  p <- shap.plot.dependence(
    shap,
    x = x,
    color_feature = "auto",

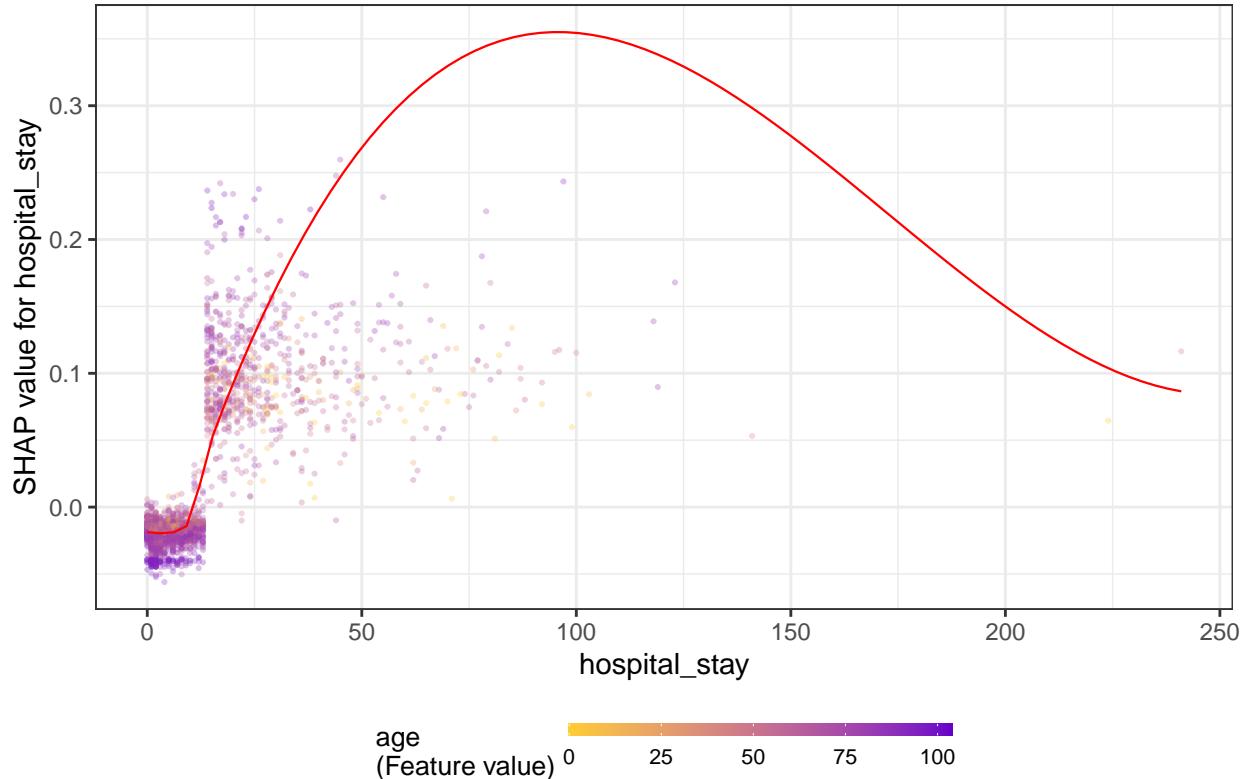
```

```

    smooth = TRUE,
    jitter_width = 0.01,
    alpha = 0.3
) +
  labs(title = x)
print(p)
}

## `geom_smooth()` using formula 'y ~ x'
```

### hospital\_stay

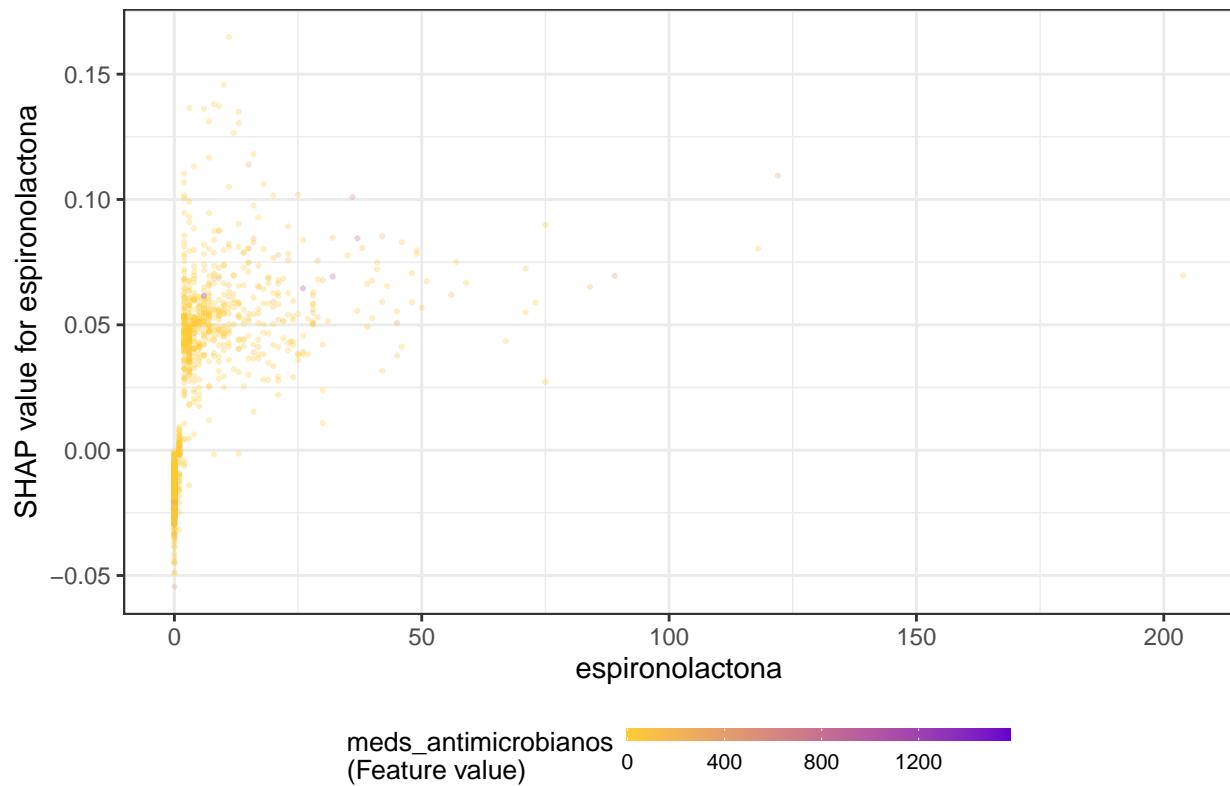


```

## `geom_smooth()` using formula 'y ~ x'

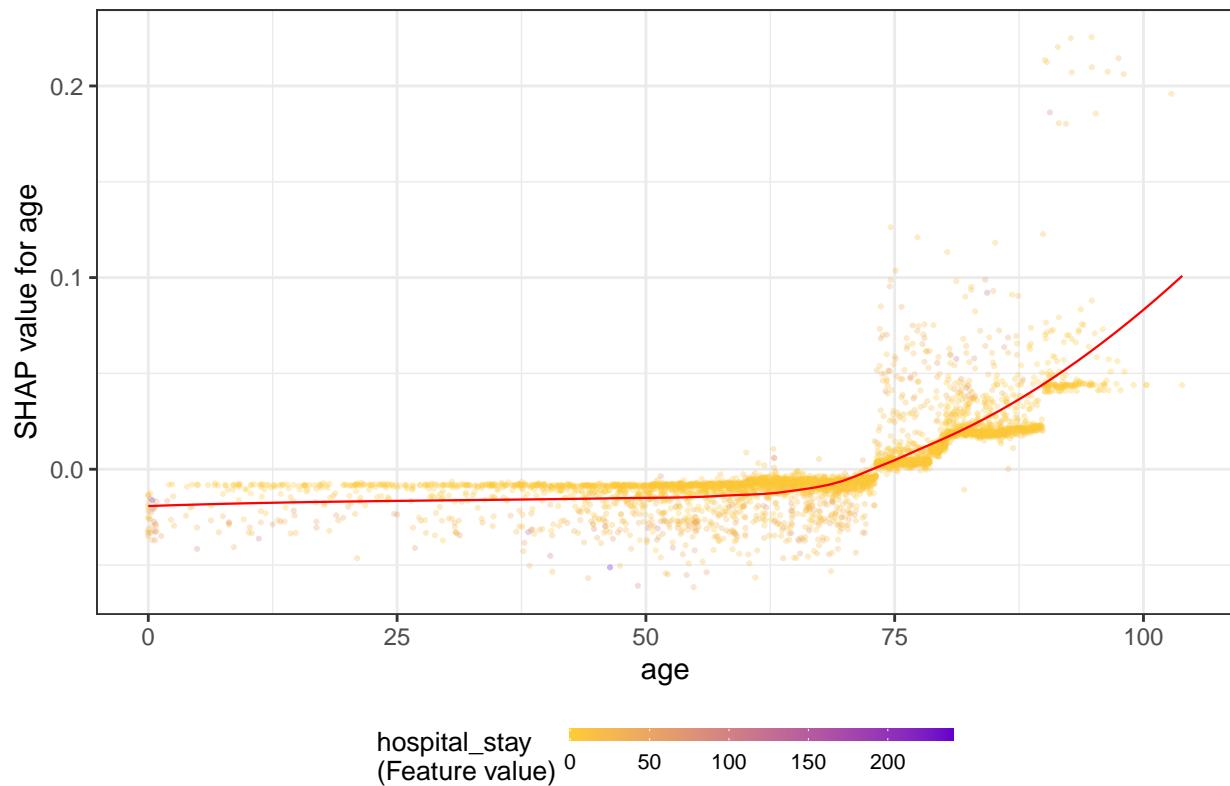
## Warning: Removed 1041 rows containing non-finite values (stat_smooth).
## Warning in simpleLoess(y, x, w, span, degree = degree, parametric = parametric, : at -1.02
## Warning in simpleLoess(y, x, w, span, degree = degree, parametric = parametric, : radius 1.0404
## Warning in simpleLoess(y, x, w, span, degree = degree, parametric = parametric, : all data on boundary
## of neighborhood. make span bigger
## Warning in simpleLoess(y, x, w, span, degree = degree, parametric = parametric, : pseudoinverse used at
## -1.02
## Warning in simpleLoess(y, x, w, span, degree = degree, parametric = parametric, : neighborhood radius
## 1.02
## Warning in simpleLoess(y, x, w, span, degree = degree, parametric = parametric, : reciprocal condition
## number 1
## Warning in simpleLoess(y, x, w, span, degree = degree, parametric = parametric, : zero-width
## neighborhood. make span bigger
## Warning: Computation failed in 'stat_smooth()':
## NA/NaN/Inf in foreign function call (arg 5)
## Warning: Removed 1041 rows containing missing values (geom_point).
```

### espironolactona



```
## `geom_smooth()` using formula 'y ~ x'
```

### age

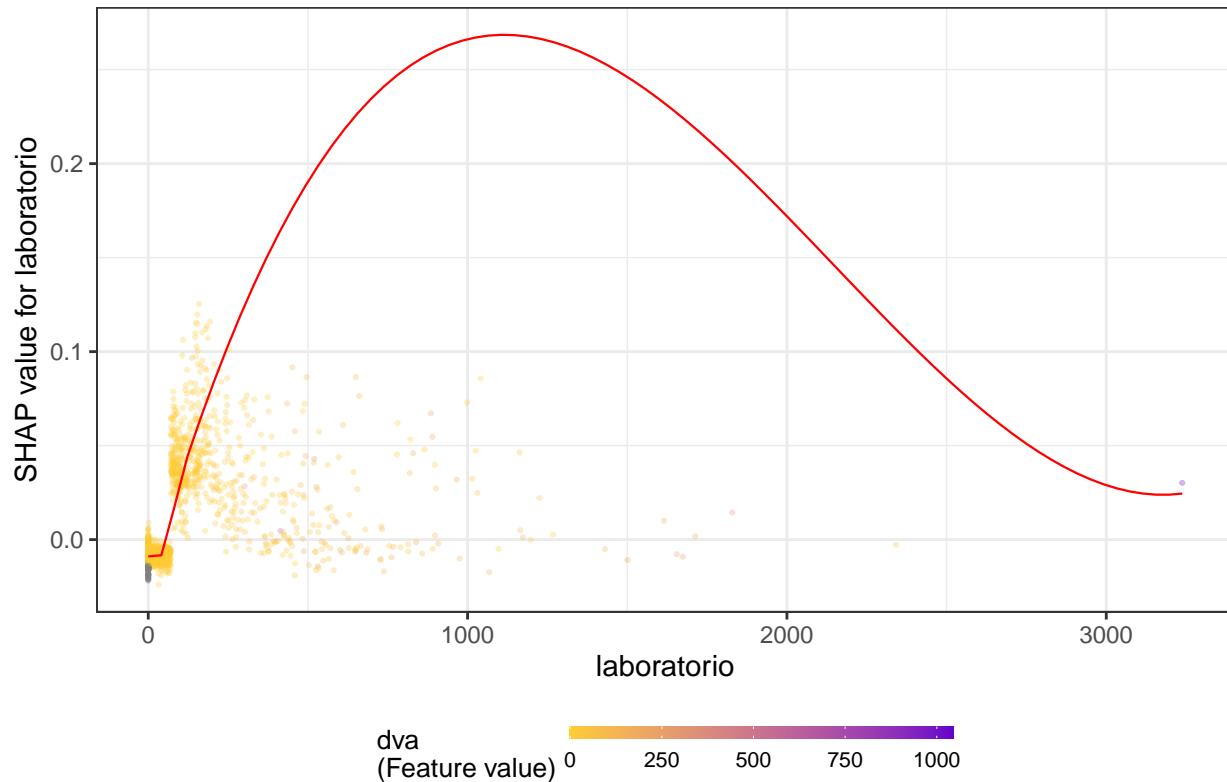


```
## `geom_smooth()` using formula 'y ~ x'
```

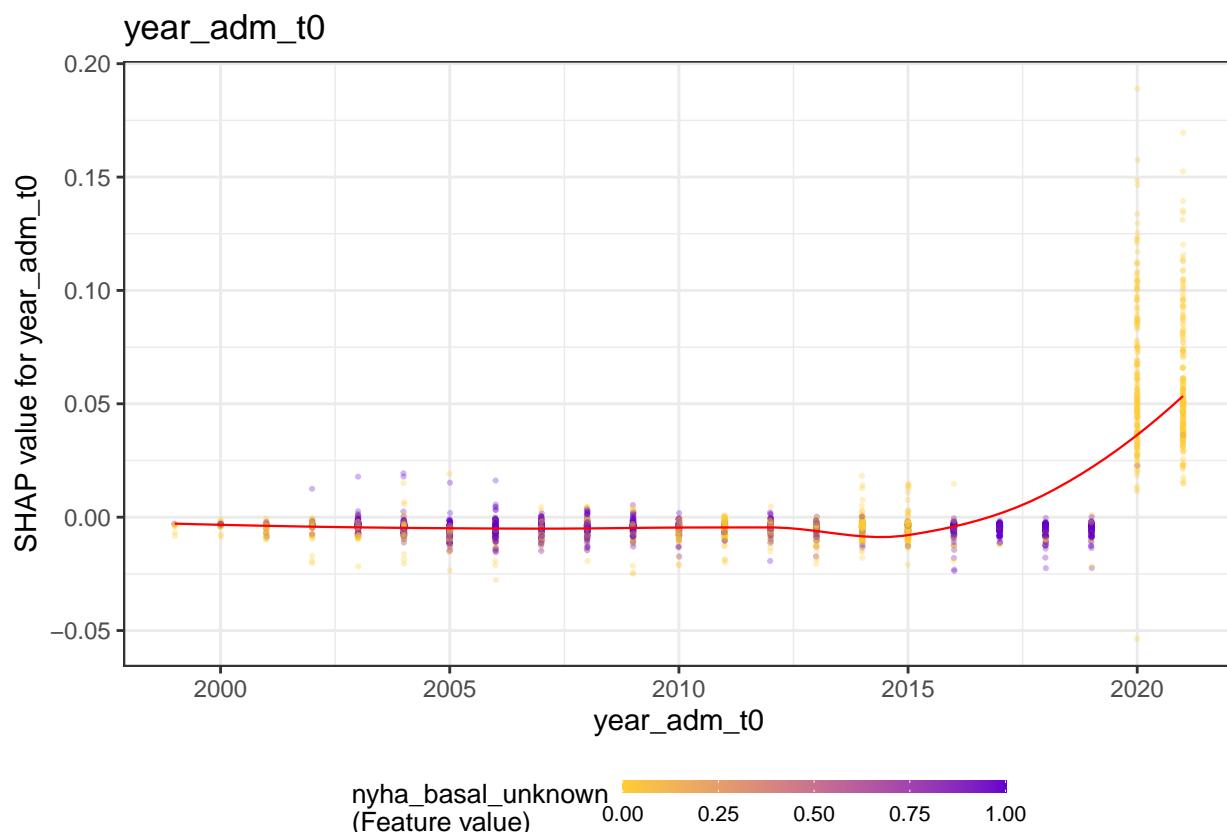
```
## Warning: Removed 808 rows containing non-finite values (stat_smooth).
```

```
## Warning: Removed 808 rows containing missing values (geom_point).
```

## laboratorio



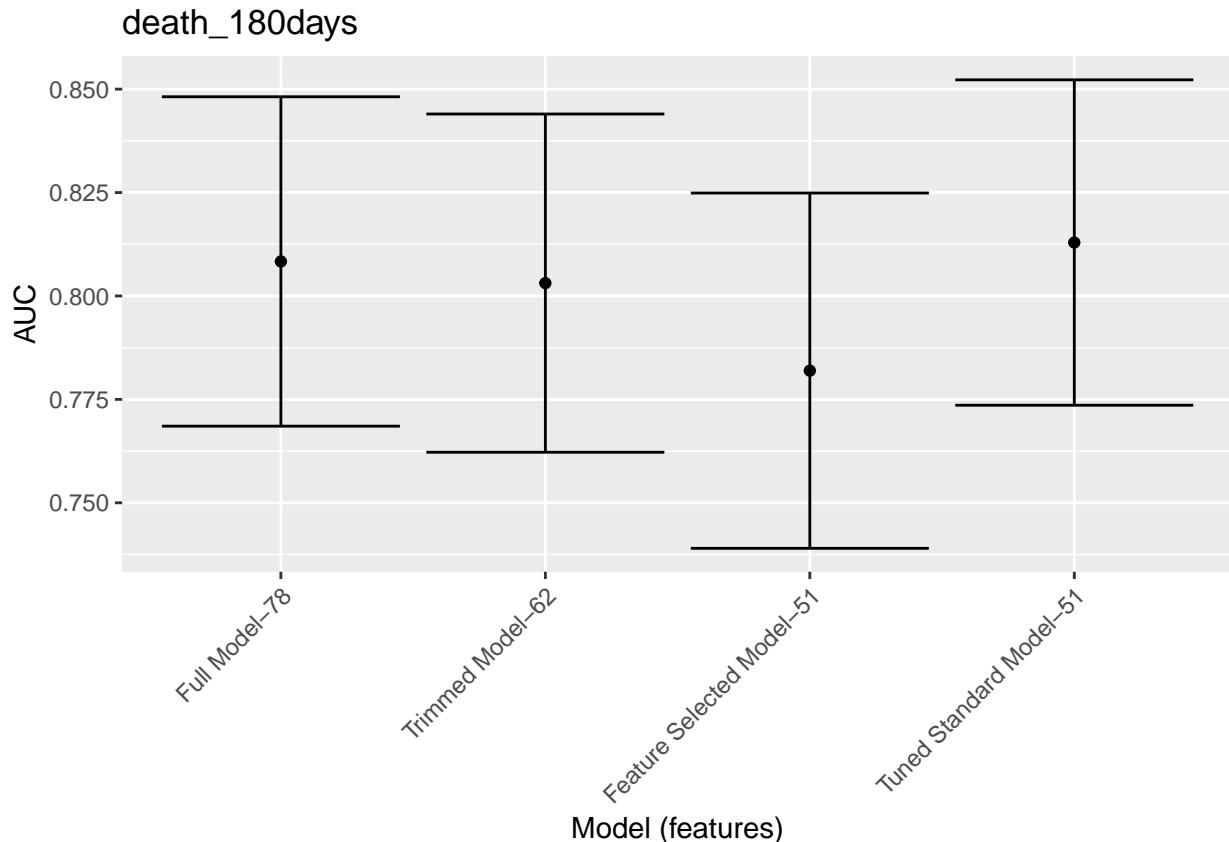
```
## `geom_smooth()` using formula 'y ~ x'  
## Warning: Removed 7 rows containing non-finite values (stat_smooth).  
## Warning: Removed 7 rows containing missing values (geom_point).
```



## Models Comparison

```
df_auc <- tibble::tribble(
  ~Model, `~AUC`, `~Lower Limit`, `~Upper Limit`, `~Features`,
  'Full Model', full_model$auc, full_model$auc_lower, full_model$auc_upper, length(features),
  'Trimmed Model', trimmed_model$auc, trimmed_model$auc_lower, trimmed_model$auc_upper, length(trimmed_features),
  'Feature Selected Model', feature_selected_model$auc, feature_selected_model$auc_lower, feature_selected_model$auc_upper, length(selected_features),
  'Tuned Standard Model', standard_results$auc, standard_results$auc_lower, standard_results$auc_upper, length(selected_features),
  # 'Tuned Smote Model', smote_results$auc, smote_results$auc_lower, smote_results$auc_upper, length(selected_features),
  # 'Tuned Upsample Model', upsample_results$auc, upsample_results$auc_lower, upsample_results$auc_upper, length(selected_features)
) %>%
  mutate(Target = outcome_column,
    `Model (features)` = fct_reorder(paste0(Model, " - ", Features), -Features))

df_auc %>%
  ggplot(aes(
    x = `Model (features)``,
    y = AUC,
    ymin = `Lower Limit``,
    ymax = `Upper Limit``
  )) +
  geom_point() +
  geom_errorbar() +
  labs(title = outcome_column) +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```



```
saveRDS(df_auc, sprintf("./auxiliar/final_model/performance/%s.RData", outcome_column))
```