

Model Selection - readmission_1year

Eduardo Yuki Yada

Global parameters

```
k <- params$k # Number of folds for cross validation
grid_size <- params$grid_size # Number of parameter combination to tune on each model
repeats <- params$repeats
RUN_ALL_MODELS <- params$RUN_ALL_MODELS
Hmisc::list.tree(params)
```

```
## params = list 5 (968 bytes)
## . outcome_column = character 1= readmission_1year
## . k = double 1= 10
## . grid_size = double 1= 20
## . repeats = double 1= 2
## . RUN_ALL_MODELS = logical 1= TRUE
```

Minutes to run: 0

Imports

```
library(tidyverse)
library(yaml)
library(tidymodels)
library(usemodels)
library(vip)
library(bonsai)
library(lightgbm)
library(caret)
library(pROC)

source("aux_functions.R")
predict <- stats::predict
```

Minutes to run: 0

Loading data

```
load('dataset/processed_data.RData')
load('dataset/processed_dictionary.RData')

columns_list <- yaml.load_file("./auxiliar/columns_list.yaml")

outcome_column <- params$outcome_column
features_list <- params$features_list

df <- mutate(df, across(where(is.character), as.factor))
```

Minutes to run: 0.006

```
dir.create(file.path("./auxiliar/model_selection/hyperparameters/"),
           showWarnings = FALSE,
```

```

        recursive = TRUE)

dir.create(file.path("../auxiliar/model_selection/performance/"),
          showWarnings = FALSE,
          recursive = TRUE)

```

Minutes to run: 0

Eligible features

```

cat_features_list = read_yaml(sprintf(
  "../auxiliar/significant_columns/categorical_%s.yaml",
  outcome_column
))

num_features_list = read_yaml(sprintf(
  "../auxiliar/significant_columns/numerical_%s.yaml",
  outcome_column
))

features_list = c(cat_features_list, num_features_list)

```

Minutes to run: 0

```

eligible_columns = df_names %>%
  filter(momento.aquisicao == 'Admissão t0') %>%
  .$variable.name

exception_columns = c('death_intraop', 'death_intraop_1', 'disch_outcomes_t0')

correlated_columns = c('year_procedure_1', # com year_adm_t0
  'age_surgery_1', # com age
  'admission_t0', # com admission_pre_t0_count
  'atb', # com meds_antimicrobianos
  'classe_meds_cardio_qtde', # com classe_meds_qtde
  'suporte_hemod', # com proced_invasivos_qtde,
  'radiografia', # com exames_imagem_qtde
  'ecg' # com metodos_graficos_qtde
)

eligible_features = eligible_columns %>%
  base::intersect(c(columns_list$categorical_columns, columns_list$numerical_columns)) %>%
  setdiff(c(exception_columns, correlated_columns))

features = base::intersect(eligible_features, features_list)

gluedown::md_order(features, seq = TRUE, pad = TRUE)

## 01. sex
## 02. age
## 03. race
## 04. education_level
## 05. patient_state
## 06. underlying_heart_disease
## 07. heart_disease
## 08. nyha_basal
## 09. prior_mi
## 10. heart_failure
## 11. af
## 12. cardiac_arrest
## 13. transplant

```

14. valvopathy
15. endocardites
16. diabetes
17. renal_failure
18. hemodialysis
19. copd
20. comorbidities_count
21. procedure_type_1
22. reop_type_1
23. procedure_type_new
24. cied_final_1
25. cied_final_group_1
26. admission_pre_t0_count
27. admission_pre_t0_180d
28. year_adm_t0
29. icu_t0
30. dialysis_t0
31. admission_t0_emergency
32. aco
33. antiarritmico
34. betabloqueador
35. ieca_bra
36. dva
37. digoxina
38. estatina
39. diuretico
40. vasodilatador
41. insuf_cardiaca
42. espirolactona
43. bloq_calcio
44. antiplaquetario_ev
45. insulina
46. anticonvulsivante
47. psicofarmacos
48. antifungico
49. antiviral
50. antiretroviral
51. classe_meds_qtde
52. meds_cardiovasc_qtde
53. meds_antimicrobianos
54. ventilacao_mecanica
55. cec
56. transplante_cardiaco
57. cir_toracica
58. outros_proced_cirurgicos
59. icp
60. intervencao_cv
61. angioplastia
62. cateterismo
63. eletrofisiologia
64. cateter_venoso_central
65. proced_invasivos_qtde
66. cve_desf
67. transfusao
68. interconsulta
69. equipe_multiprof
70. holter
71. teste_esforco
72. espiro_ergoespiro
73. tilt_teste
74. metodos_graficos_qtde

```
## 75. laboratorio
## 76. cultura
## 77. analises_clinicas_qtde
## 78. citologia
## 79. biopsia
## 80. histopatologia_qtde
## 81. angio_rm
## 82. angio_tc
## 83. aortografia
## 84. arteriografia
## 85. cintilografia
## 86. ecocardiograma
## 87. endoscopia
## 88. flebografia
## 89. pet_ct
## 90. ultrassom
## 91. tomografia
## 92. ressonancia
## 93. exames_imagem_qtde
## 94. dieta_parenteral
## 95. bic
## 96. mpp
## 97. hospital_stay
```

Minutes to run: 0

Train test split (70%/30%)

```
set.seed(42)

if (outcome_column == 'readmission_30d') {
  df_split <- readRDS("./dataset/split_object.rds")
} else {
  df_split <- initial_split(df, prop = .7, strata = all_of(outcome_column))
}

df_train <- training(df_split) %>% dplyr::select(all_of(c(features, outcome_column)))
df_test <- testing(df_split) %>% dplyr::select(all_of(c(features, outcome_column)))

df_folds <- vfold_cv(df_train, v = k,
                     strata = all_of(outcome_column))
```

Minutes to run: 0.001

Boosted Tree (XGBoost)

```
xgboost_recipe <-
  recipe(formula = sprintf("%s ~ .", outcome_column) %>% as.formula, data = df_train) %>%
  step_novel(all_nominal_predictors()) %>%
  step_unknown(all_nominal_predictors()) %>%
  step_other(all_nominal_predictors(), threshold = 0.05, other = ".merged") %>%
  step_dummy(all_nominal_predictors())

xgboost_spec <- boost_tree(
  trees = tune(),
  min_n = tune(),
  tree_depth = tune(),
  learn_rate = tune(),
) %>%
  set_engine("xgboost",
```

```

      nthread = 8) %>%
set_mode("classification")

xgboost_grid <- grid_latin_hypercube(
  trees(range = c(25L, 150L)),
  min_n(range = c(2L, 100L)),
  tree_depth(range = c(2L, 15L)),
  learn_rate(range = c(-3, -1), trans = log10_trans()),
  size = grid_size
)

xgboost_workflow <-
  workflow() %>%
  add_recipe(xgboost_recipe) %>%
  add_model(xgboost_spec)

xgboost_tune <-
  xgboost_workflow %>%
  tune_grid(resamples = df_folds,
            grid = xgboost_grid)


xgboost_tune %>%
  show_best("roc_auc")

## # A tibble: 5 x 10
##   trees min_n tree_depth learn_rate .metric .estimator mean      n std_err .config
##   <int> <int>    <int>    <dbl> <chr>   <chr>    <dbl> <int>  <dbl> <chr>
## 1    96    92        5    0.0960 roc_auc binary    0.718    10 0.00757 Prepro~
## 2   136    81        2    0.0413 roc_auc binary    0.716    10 0.00727 Prepro~
## 3   130    29        4    0.0226 roc_auc binary    0.716    10 0.00681 Prepro~
## 4   140    43       14    0.0649 roc_auc binary    0.715    10 0.00767 Prepro~
## 5    46    33       13    0.0524 roc_auc binary    0.711    10 0.00731 Prepro~

best_xgboost <- xgboost_tune %>%
  select_best("roc_auc")

autoplot(xgboost_tune, metric = "roc_auc")

```

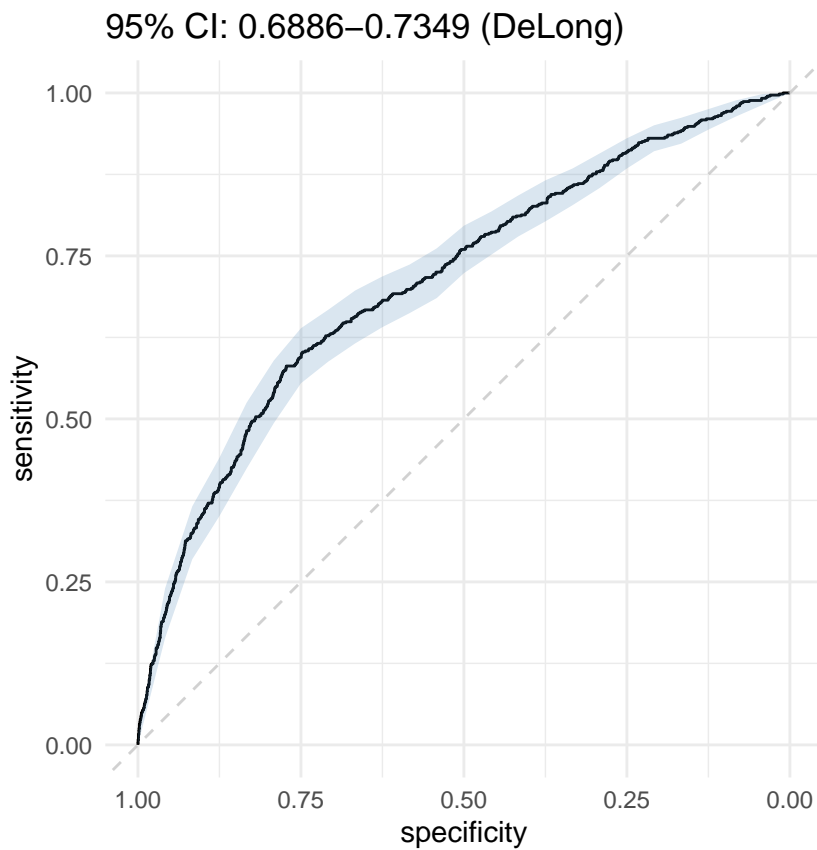


```
final_xgboost_workflow <-
  xgboost_workflow %>%
  finalize_workflow(best_xgboost)

last_xgboost_fit <-
  final_xgboost_workflow %>%
  last_fit(df_split)


final_xgboost_fit <- extract_workflow(last_xgboost_fit)

xgboost_auc <- validation(final_xgboost_fit, df_test)
```



```
## [1] "Optimal Threshold: 0.14"
## Confusion Matrix and Statistics
##
##      reference
## data    0    1
##    0 3188 253
##    1  939 351
##
##              Accuracy : 0.748
##              95% CI   : (0.7354, 0.7604)
##    No Information Rate : 0.8723
##    P-Value [Acc > NIR] : 1
##
##              Kappa   : 0.2382
##
## Mcnemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.7725
##              Specificity : 0.5811
##              Pos Pred Value : 0.9265
##              Neg Pred Value : 0.2721
##              Prevalence : 0.8723
##              Detection Rate : 0.6739
##              Detection Prevalence : 0.7273
##              Balanced Accuracy : 0.6768
##
##              'Positive' Class : 0
##
```

```
extract_vip(final_xgboost_fit, pred_wrapper = predict,
            reference_class = "0")
```



```
xgboost_parameters <- xgboost_tune %>%
  show_best("roc_auc", n = 1) %>%
  select(-.metric, -.estimator, -.config, -mean, -n, -std_err) %>%
  as.list
```

Minutes to run: 5.452

Boosted Tree (LightGBM)

```
lightgbm_recipe <-
  recipe(formula = sprintf("%s ~ .", outcome_column) %>% as.formula, data = df_train) %>%
  step_novel(all_nominal_predictors()) %>%
  step_unknown(all_nominal_predictors()) %>%
  step_other(all_nominal_predictors(), threshold = 0.05, other = ".merged") %>%
  step_dummy(all_nominal_predictors())

lightgbm_spec <- boost_tree(
  trees = tune(),
  min_n = tune(),
  tree_depth = tune(),
  learn_rate = tune(),
  sample_size = 1
) %>%
  set_engine("lightgbm",
    nthread = 8) %>%
  set_mode("classification")

lightgbm_grid <- grid_latin_hypercube(
  trees(range = c(25L, 150L)),
  min_n(range = c(2L, 100L)),
  tree_depth(range = c(2L, 15L)),
  learn_rate(range = c(-3, -1), trans = log10_trans()),
  size = grid_size
```



```
)
```

```
lightgbm_workflow <-  
  workflow() %>%  
  add_recipe(lightgbm_recipe) %>%  
  add_model(lightgbm_spec)
```


```
lightgbm_tune <-  
  lightgbm_workflow %>%  
  tune_grid(resamples = df_folds,  
            grid = lightgbm_grid)
```

```
lightgbm_tune %>%  
  show_best("roc_auc")
```

```
## # A tibble: 5 x 10  
##   trees min_n tree_depth learn_rate .metric .estimator mean      n std_err .config  
##   <int> <int>    <int>    <dbl> <chr>   <chr>    <dbl> <int>  <dbl> <chr>  
## 1   143     6      10    0.0227 roc_auc binary    0.721    10 0.00741 Prepro~  
## 2    65    16       3    0.0417 roc_auc binary    0.720    10 0.00666 Prepro~  
## 3   114   100       4    0.0575 roc_auc binary    0.719    10 0.00703 Prepro~  
## 4   135    26      14    0.0180 roc_auc binary    0.718    10 0.00751 Prepro~  
## 5    42    72       6    0.0951 roc_auc binary    0.717    10 0.00744 Prepro~
```

```
best_lightgbm <- lightgbm_tune %>%  
  select_best("roc_auc")
```

```
autoplot(lightgbm_tune, metric = "roc_auc")
```



```
final_lightgbm_workflow <-  
  lightgbm_workflow %>%  
  finalize_workflow(best_lightgbm)
```

```
last_lightgbm_fit <-
```


```

final_lightgbm_workflow %>%
last_fit(df_split)

final_lightgbm_fit <- extract_workflow(last_lightgbm_fit)

lightgbm_auc <- validation(final_lightgbm_fit, df_test)

```



```

## [1] "Optimal Threshold: 0.11"
## Confusion Matrix and Statistics
##
##      reference
## data    0    1
##    0 2804  211
##    1 1323  393
##
##              Accuracy : 0.6758
##              95% CI : (0.6622, 0.6891)
##    No Information Rate : 0.8723
##    P-Value [Acc > NIR] : 1
##
##              Kappa : 0.1848
##
## Mcnemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.6794
##              Specificity : 0.6507
##              Pos Pred Value : 0.9300
##              Neg Pred Value : 0.2290
##              Prevalence : 0.8723
##              Detection Rate : 0.5927
##              Detection Prevalence : 0.6373
##              Balanced Accuracy : 0.6650
##

```

```
##      'Positive' Class : 0
##
```

```
lightgbm_parameters <- lightgbm_tune %>%
  show_best("roc_auc", n = 1) %>%
  select(-.metric, -.estimator, -.config, -mean, -n, -std_err) %>%
  as.list
```


```
Hmisc::list.tree(lightgbm_parameters)
```

```
## lightgbm_parameters = list 4 (736 bytes)
## . trees = integer 1= 143
## . min_n = integer 1= 6
## . tree_depth = integer 1= 10
## . learn_rate = double 1= 0.022707
```

```
con <- file(sprintf('./auxiliar/model_selection/hyperparameters/%s.yaml', outcome_column), "w")
write_yaml(lightgbm_parameters, con)
close(con)
```

Minutes to run: 3.568

Histogram of preds



Minutes to run:

0.017

GLM

```
glmnet_recipe <-
  recipe(formula = sprintf("%s ~ .", outcome_column) %>% as.formula, data = df_train) %>%
  step_novel(all_nominal_predictors()) %>%
  step_unknown(all_nominal_predictors()) %>%
  step_other(all_nominal_predictors(), threshold = 0.05, other = ".merged") %>%
  step_dummy(all_nominal_predictors()) %>%
  step_zv(all_predictors()) %>%
  step_normalize(all_numeric_predictors())
```

```


glmnet_spec <-
  logistic_reg(penalty = 0) %>%
  set_mode("classification") %>%
  set_engine("glm")

glmnet_workflow <-
  workflow() %>%
  add_recipe(glmnet_recipe) %>%
  add_model(glmnet_spec)

glm_fit <- glmnet_workflow %>%
  fit(df_train)

glmnet_auc <- validation(glm_fit, df_test)

```



```


## [1] "Optimal Threshold: 0.11"
## Confusion Matrix and Statistics
##
##      reference
## data    0    1
## 0 1709  136
## 1   811  243
##
##              Accuracy : 0.6733
##              95% CI   : (0.6559, 0.6904)
##    No Information Rate : 0.8693
##    P-Value [Acc > NIR] : 1
##
##              Kappa   : 0.1818
##
##  McNemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.6782

```


```
##          Specificity : 0.6412
##          Pos Pred Value : 0.9263
##          Neg Pred Value : 0.2306
##          Prevalence : 0.8693
##          Detection Rate : 0.5895
##          Detection Prevalence : 0.6364
##          Balanced Accuracy : 0.6597
##
##          'Positive' Class : 0
##
```

```
pfun_glmnet <- function(object, newdata) predict(object, newx = newdata)

extract_vip(glm_fit, pred_wrapper = pfun_glmnet,
            reference_class = "1", method = 'model')
```



```
extract_vip(glm_fit, pred_wrapper = pfun_glmnet,
            reference_class = "1", method = 'permute')
```



Importance

Minutes to run:

0.206


Decision Tree

```
tree_recipe <-  
  recipe(formula = sprintf("%s ~ .", outcome_column) %>% as.formula, data = df_train) %>%  
  step_novel(all_nominal_predictors()) %>%  
  step_unknown(all_nominal_predictors()) %>%  
  step_other(all_nominal_predictors(), threshold = 0.05, other = ".merged") %>%  
  step_dummy(all_nominal_predictors()) %>%  
  step_zv(all_predictors())  
  
tree_spec <-  
  decision_tree(cost_complexity = tune(),  
                tree_depth = tune(),  
                min_n = tune()) %>%  
  set_mode("classification") %>%  
  set_engine("rpart")  
  
tree_grid <- grid_latin_hypercube(cost_complexity(),  
                                  tree_depth(),  
                                  min_n(),  
                                  size = grid_size)  
  
tree_workflow <-  
  workflow() %>%  
  add_recipe(tree_recipe) %>%  
  add_model(tree_spec)  
  
tree_tune <-  
  tree_workflow %>%  
  tune_grid(resamples = df_folds,  
            grid = tree_grid)
```

```
tree_tune %>%
  collect_metrics()
```

```
## # A tibble: 40 x 9
##   cost_complexity tree_depth min_n .metric .estimator mean      n std_err .config
##   <dbl>          <int> <int> <chr>  <chr>    <dbl> <int>  <dbl> <chr>
## 1      2.57e-10         7    17 accura~ binary  0.866   10 1.56e-3 Prepro~
## 2      2.57e-10         7    17 roc_auc binary  0.682   10 5.49e-3 Prepro~
## 3      7.42e- 5         4    15 accura~ binary  0.871   10 8.72e-4 Prepro~
## 4      7.42e- 5         4    15 roc_auc binary  0.661   10 7.02e-3 Prepro~
## 5      3.42e- 3        10    21 accura~ binary  0.870   10 1.14e-3 Prepro~
## 6      3.42e- 3        10    21 roc_auc binary  0.631   10 4.48e-3 Prepro~
## 7      3.99e- 5        14     7 accura~ binary  0.840   10 2.67e-3 Prepro~
## 8      3.99e- 5        14     7 roc_auc binary  0.605   10 7.61e-3 Prepro~
## 9      1.89e- 5        10     6 accura~ binary  0.854   10 1.49e-3 Prepro~
## 10     1.89e- 5        10     6 roc_auc binary  0.644   10 8.08e-3 Prepro~
## # i 30 more rows
```

```
autoplot(tree_tune, metric = "roc_auc")
```



```
tree_tune %>%
  show_best("roc_auc")
```

```
## # A tibble: 5 x 9
##   cost_complexity tree_depth min_n .metric .estimator mean      n std_err .config
##   <dbl>          <int> <int> <chr>  <chr>    <dbl> <int>  <dbl> <chr>
## 1      3.97e- 9         6    23 roc_auc binary  0.686   10 0.00538 Preproc~
## 2      2.09e- 9         6    10 roc_auc binary  0.685   10 0.00542 Preproc~
## 3      2.57e-10         7    17 roc_auc binary  0.682   10 0.00549 Preproc~
## 4      4.00e- 6         8    33 roc_auc binary  0.681   10 0.00453 Preproc~
## 5      3.30e- 8         5    25 roc_auc binary  0.678   10 0.00769 Preproc~
```

```
best_tree <- tree_tune %>%
  select_best("roc_auc")
```

```


final_tree_workflow <-
  tree_workflow %>%
  finalize_workflow(best_tree)

last_tree_fit <-
  final_tree_workflow %>%
  last_fit(df_split)

final_tree_fit <- extract_workflow(last_tree_fit)

tree_auc <- validation(final_tree_fit, df_test)

```



```

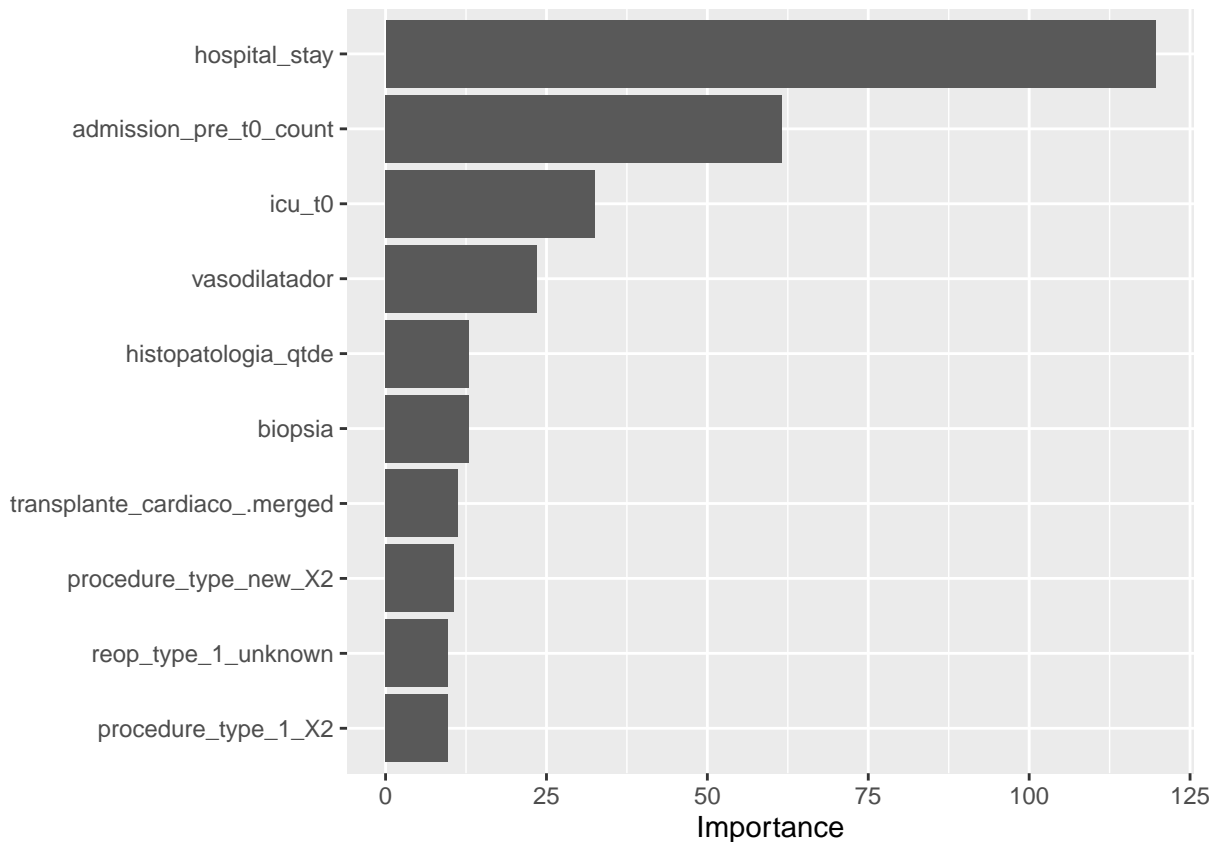
## [1] "Optimal Threshold: 0.11"
## Confusion Matrix and Statistics
##
##      reference
## data    0    1
##    0 2946  257
##    1 1181  347
##
##              Accuracy : 0.696
##              95% CI : (0.6827, 0.7091)
##    No Information Rate : 0.8723
##    P-Value [Acc > NIR] : 1
##
##              Kappa : 0.1744
##
##  McNemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.7138
##              Specificity : 0.5745
##              Pos Pred Value : 0.9198
##              Neg Pred Value : 0.2271

```



```
##           Prevalence : 0.8723
##           Detection Rate : 0.6227
##           Detection Prevalence : 0.6770
##           Balanced Accuracy : 0.6442
##
##           'Positive' Class : 0
##
```

```
extract_vip(final_tree_fit, pred_wrapper = predict,
             reference_class = "0", use_matrix = FALSE,
             method = 'model')
```



```
# extract_vip(final_tree_fit, pred_wrapper = predict,
#             reference_class = "1", use_matrix = FALSE,
#             method = 'permute')
```

Minutes to run: 4.95

Random Forest

```
rf_recipe <-
  recipe(formula = sprintf("%s ~ .", outcome_column) %>% as.formula,
          data = df_train) %>%
  step_novel(all_nominal_predictors()) %>%
  step_unknown(all_nominal_predictors()) %>%
  step_other(all_nominal_predictors(), threshold = 0.05, other = ".merged") %>%
  step_dummy(all_nominal_predictors()) %>%
  step_zv(all_predictors()) %>%
  step_impute_mean(all_numeric_predictors())

rf_spec <-
  rand_forest(mtry = tune(),
              trees = tune(),
              min_n = tune()) %>%
```

```

set_mode("classification") %>%
set_engine("randomForest",
  probability = TRUE,
  nthread = 8)

rf_grid <- grid_latin_hypercube(mtry(range = c(1L, 50L)),
  trees(range = c(100L, 300L)),
  min_n(),
  size = grid_size)

rf_workflow <-
  workflow() %>%
  add_recipe(rf_recipe) %>%
  add_model(rf_spec)


rf_tune <-
  rf_workflow %>%
  tune_grid(resamples = df_folds,
    grid = rf_grid)

rf_tune %>%
  collect_metrics()

## # A tibble: 40 x 9
##   mtry trees min_n .metric .estimator mean      n std_err .config
##   <int> <int> <int> <chr>   <chr>   <dbl> <int>   <dbl> <chr>
## 1     2    111    10 accuracy binary    0.873    10 0.000406 Preprocessor1_Model~
## 2     2    111    10 roc_auc  binary    0.671    10 0.00754  Preprocessor1_Model~
## 3     5    210    23 accuracy binary    0.874    10 0.000752 Preprocessor1_Model~
## 4     5    210    23 roc_auc  binary    0.700    10 0.00692  Preprocessor1_Model~
## 5    35    225     5 accuracy binary    0.874    10 0.00138  Preprocessor1_Model~
## 6    35    225     5 roc_auc  binary    0.698    10 0.00806  Preprocessor1_Model~
## 7    19    264    15 accuracy binary    0.875    10 0.00122  Preprocessor1_Model~
## 8    19    264    15 roc_auc  binary    0.707    10 0.00787  Preprocessor1_Model~
## 9    42    168    29 accuracy binary    0.874    10 0.00110  Preprocessor1_Model~
## 10   42    168    29 roc_auc  binary    0.703    10 0.00781  Preprocessor1_Model~
## # i 30 more rows

autoplot(rf_tune, metric = "roc_auc")

```



```
rf_tune %>%
  show_best("roc_auc")
```

```
## # A tibble: 5 x 9
##   mtry trees min_n .metric .estimator  mean     n std_err .config
##   <int> <int> <int> <chr>   <chr>    <dbl> <int>   <dbl> <chr>
## 1    15   154    13 roc_auc binary    0.710    10 0.00870 Preprocessor1_Model117
## 2     9    90    17 roc_auc binary    0.707    10 0.00882 Preprocessor1_Model108
## 3    19   264    15 roc_auc binary    0.707    10 0.00787 Preprocessor1_Model104
## 4    29   257    34 roc_auc binary    0.705    10 0.00774 Preprocessor1_Model116
## 5    21   243    35 roc_auc binary    0.705    10 0.00796 Preprocessor1_Model106
```


```
best_rf <- rf_tune %>%
  select_best("roc_auc")
```

```
final_rf_workflow <-
  rf_workflow %>%
  finalize_workflow(best_rf)
```

```
last_rf_fit <-
  final_rf_workflow %>%
  last_fit(df_split)
```

```
final_rf_fit <- extract_workflow(last_rf_fit)
```


```
rf_auc <- validation(final_rf_fit, df_test)
```



```
## [1] "Optimal Threshold: 0.09"
## Confusion Matrix and Statistics
##
##      reference
## data    0    1
##    0 2974  251
##    1 1153  353
##
##              Accuracy : 0.7032
##              95% CI   : (0.69, 0.7162)
##    No Information Rate : 0.8723
##    P-Value [Acc > NIR] : 1
##
##              Kappa : 0.1863
##
##  McNemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.7206
##              Specificity : 0.5844
##    Pos Pred Value : 0.9222
##    Neg Pred Value : 0.2344
##    Prevalence : 0.8723
##    Detection Rate : 0.6286
##    Detection Prevalence : 0.6817
##    Balanced Accuracy : 0.6525
##
##    'Positive' Class : 0
##
```

```
pfun_rf <- function(object, newdata) predict(object, data = newdata)

extract_vip(final_rf_fit, pred_wrapper = predict,
  reference_class = "1", use_matrix = FALSE,
  method = 'model')
```




```
# extract_vip(final_rf_fit, pred_wrapper = predict,
#             reference_class = "1", use_matrix = FALSE,
#             method = 'permute')
```

Minutes to run: 67.47

Models Comparison

```
if (RUN_ALL_MODELS) {
  df_auc <- tibble::tribble(
    ~Model, ~`AUC`, ~`Lower Limit`, ~`Upper Limit`,
    'XGBoost', as.numeric(xgboost_auc$auc), xgboost_auc$ci[1], xgboost_auc$ci[3],
    'LightGBM', as.numeric(lightgbm_auc$auc), lightgbm_auc$ci[1], lightgbm_auc$ci[3],
    'GLM', as.numeric(glmnet_auc$auc), glmnet_auc$ci[1], glmnet_auc$ci[3],
    'Decision Tree', as.numeric(tree_auc$auc), tree_auc$ci[1], tree_auc$ci[3],
    'Random Forest', as.numeric(rf_auc$auc), rf_auc$ci[1], rf_auc$ci[3]
  ) %>%
    mutate(Target = outcome_column)
} else {
  df_auc <- tibble::tribble(
    ~Model, ~`AUC`, ~`Lower Limit`, ~`Upper Limit`,
    'LightGBM', as.numeric(lightgbm_auc$auc), lightgbm_auc$ci[1], lightgbm_auc$ci[3]
  ) %>%
    mutate(Target = outcome_column)
}

df_auc %>%
  ggplot(aes(x = Model, y = AUC, ymin = `Lower Limit`, ymax = `Upper Limit`)) +
  geom_point() +
  geom_errorbar()
```



```
write_csv(df_auc, sprintf("./auxiliar/model_selection/performance/%s.csv", outcome_column))
```

Minutes to run: 0.002